

# **Introduction to High Performance Scientific Computing**

**Autumn, 2016**

**Lecture 10**

# Today

---

- **A few comments on Fortran**
- **F2Py with fortran modules**
- **Timing Fortran code**
- **Networks**

# Notes on Fortran

---

- **Do not attempt to guess Fortran syntax! – look it up instead**
- **Do not develop Fortran code in Python using f2py – get a Fortran-only code working first**
- **Test code (compile and run) after adding a few (4-5) lines of code. Don't write 100 lines and then test!**
- **Lecture 7: code structure, variable types, loops, if-then, subroutines**
- **Lecture 8: allocatable arrays, functions modules**
- **Lecture 9: Lapack, f2py**

# Notes on Fortran

---

## Debugging code

- If the compiler is giving you a series of errors, look at the topmost error message. It will include a line number
- If the code runs but crashes with a segmentation fault: probably a problem with array indices (e.g. trying to access the 12<sup>th</sup> element of a dimension(11) array)
- If you can't tell where the code is crashing, add print statements, e.g. `print *, 1` --some code-- `print *, 2`  
When the code is run, if the 1 prints to screen but not the 2, you know where the problem is.
- If the code runs, but gives the wrong answer, add print statements outputting values of variables, try a small problem size where you know what values the variables can take

# Notes on Fortran

---

- Modules consist of:
    1. Module variables
    2. Module sub-programs
  - Module variables are “available” throughout the module
    - They *do not* need to be declared
    - They *do not* need to be provided as input/output in the sub-program header
  - Module variables and module sub-programs are also “available” in any program or sub-program that *uses* the module
  - A module by itself doesn’t do anything
    - There should be a “main” program which uses it
    - You can compile a module by itself: `gfortran -c module.f90`
    - But to generate an executable, you need a program:  
`gfortran -o program.exe module.f90 program.f90`
- When re-compiling, first remove the previous .mod file

# Notes on Fortran

---

- **Compiling code that uses lapack routines:**

```
$ gfortran -c program.f90  
$ gfortran -o program.exe program.o -llapack  
or
```

```
$ gfortran -o program.exe program.f90 -llapack
```

**Similarly, with f2py:**

```
$ f2py -llapack -c program.f90 -m module_name
```

# F2Py and fortran modules

---

- F2Py will recognize subroutines and functions in modules
- What about variables?
  - Try f2py with circle module from last week (*f2pymodule\_circle.f90*)

```
$ f2py -c f2pymodule_circle.f90 -m cmod
```

```
In [10]: import cmod
```

```
In [11]: cmod.<tab>  
cmod.circle  cmod.so
```

Need to look at *cmod.circle*

# F2Py and fortran modules

---

Need to look at *cmod.circle*:

```
In [12]: cmod.circle?
```

**Docstring:**

```
'd'-scalar  
initialize_pi()
```

```
Wrapper for ``initialize_pi``.
```

```
circumference = circumference(radius)
```

```
Wrapper for ``circumference``.
```

Parameters

-----

```
radius : input float
```

Returns

-----

```
circumference : float  
area = area(radius)
```

**And similar info for “area”**



# F2Py and fortran modules

---

How do we access variables and methods in *cmod.circle*?

```
In [23]: cmod.circle.<tab>
cmod.circle.area          cmod.circle.circumference
cmod.circle.initialize_pi  cmod.circle.pi
```

Can initialize *pi* in python:

```
In [9]: cmod.circle.pi
Out[9]: array(0.0)

In [10]: cmod.circle.pi = pi

In [11]: cmod.circle.pi
Out[11]: array(3.141592653589793)
```

# F2Py and fortran modules

---

How do we access variables and methods in *cmmod.circle*?

```
In [23]: cmmod.circle.<tab>
cmmod.circle.area          cmmod.circle.circumference
cmmod.circle.initialize_pi  cmmod.circle.pi
```

Can initialize *pi* in python:

```
In [9]: cmmod.circle.pi
Out[9]: array(0.0)

In [10]: cmmod.circle.pi = pi

In [11]: cmmod.circle.pi
Out[11]: array(3.141592653589793)
```

Can also initialize allocatable arrays, see *f2pymodule\_circle.f90...*

# F2Py and fortran modules

---

Can also initialize allocatable arrays, see *f2pymodule\_circle\_array.f90*:

```
!module for computing circumference, area, and "mass" of circle
module circle
  implicit none
  real(kind=8) :: pi
  real(kind=8), allocatable, dimension(:) :: weights, mass
  save
```

...

...

```
subroutine compute_mass(radius, mass)
  !compute mass = weights*area
  implicit none
  real(kind=8), intent(in) :: radius
  real(kind=8), intent(out) :: mass(:)

  mass = weights*area(radius)
end subroutine compute_mass
```

# F2Py and fortran modules

---

## Basic steps:

1. Compile with f2py: `$ f2py -c f2pymodule_circle_array.f90 -m cmoda`

# F2Py and fortran modules

---

## Basic steps:

1. Compile with f2py: `$ f2py -c f2pymodule_circle_array.f90 -m cmoda`

2. Import module in python: `In [4]: import cmoda`

# F2Py and fortran modules

---

## Basic steps:

1. Compile with f2py: `$ f2py -c f2pymodule_circle_array.f90 -m cmoda`

2. Import module in python: `In [4]: import cmoda`

## 3. Initialize variables:

```
In [15]: cmoda.circle.pi=pi
```

```
In [16]: cmoda.circle.weights=arange(5)
```

```
In [17]: cmoda.circle.pi,cmoda.circle.weights
```

```
Out[17]: (array(3.141592653589793), array([ 0.,  1.,  2.,  3.,  4.]))
```

# F2Py and fortran modules

---

## Basic steps:

1. Compile with f2py: `$ f2py -c f2pymodule_circle_array.f90 -m cmoda`

2. Import module in python: `In [4]: import cmoda`

## 3. Initialize variables:

```
In [15]: cmoda.circle.pi=pi
```

```
In [16]: cmoda.circle.weights=arange(5)
```

```
In [17]: cmoda.circle.pi,cmoda.circle.weights
```

```
Out[17]: (array(3.141592653589793), array([ 0.,  1.,  2.,  3.,  4.]))
```

## 4. Compute mass:

```
In [18]: cmoda.circle.compute_mass(2.0)
```

```
In [19]: cmoda.circle.mass
```

```
Out[19]: array([ 0. , 12.56637061, 25.13274123, 37.69911184,
50.26548246])
```

# Finite difference methods

---

- Finite difference methods are a standard approach for numerical differentiation.
- They form the basis for a wide variety of methods used to solve partial differential equations
- See online supplementary class notes: *Notes on numerical differentiation with finite difference methods*



# Unix *time* command

---

Can use *time* to obtain timing info for any unix command:

```
$ time ./midpoint.exe
N=      512000
sum=    3.1415926535901515
error=   3.5837999234900053E-013

real    0m0.015s
user    0m0.010s
sys 0m0.003s
```

- **real is approximately the wall-clock time**
- **user is time spent executing the program**
- **sys is time spent on system tasks required by program**

# Fortran timing functions

---

- Unix *time* doesn't tell you how much time different parts of program take
- `system_clock` and `cpu_time` gives wall time and cpu time between two points in code
- See `midpoint_time.f90`:

```
!timing variables
real(kind=8) :: cpu_t1,cpu_t2,clock_time
integer(kind=8) :: clock_t1,clock_t2,clock_rate
```

# Fortran timing functions

---

- Unix *time* doesn't tell you how much time different parts of program take
- `system_clock` and `cpu_time` gives wall time and cpu time between two points in code
- See `midpoint_time.f90`:

```
!timing variables
  real(kind=8) :: cpu_t1,cpu_t2,clock_time
  integer(kind=8) :: clock_t1,clock_t2,clock_rate

  call system_clock(clock_t1)
  call cpu_time(cpu_t1)
  !loop over intervals computing each interval's contribution to
  integral
... Midpoint quadrature ...
  call cpu_time(cpu_t2)
  print *, 'elapsed cpu time (seconds) =',cpu_t2-cpu_t1

  call system_clock(clock_t2,clock_rate)
  print *, 'elapsed wall time (seconds)= ',
          dble(clock_t2-clock_t1)/dble(clock_rate)
```

# Fortran timing functions

---

- Unix *time* doesn't tell you how much time different parts of program take
- `system_clock` and `cpu_time` gives wall time and cpu time between two points in code
- See `midpoint_time.f90`:  

```
$ ./midpoint_t.exe  
elapsed cpu time (seconds) = 8.6239999999999997E-003  
elapsed wall time (seconds)= 9.12799966E-03  
N= 512000  
sum= 3.1415926535901515  
error= 3.5837999234900053E-013
```
- Can place timing commands throughout code to find bottlenecks

# Fortran timing functions

---

- Also, often have a *theoretical estimate* of how cost scales with problem size
- A method may require  $O(N)$  (or  $O(N \ln_2 N)$  or  $O(N^2)$ ) operations
- But does your implementation of the algorithm match theory?
- How do compiler optimizations affect performance?
- Carefully timing code while varying the problem size can help answer these questions

# Networks

---

**Examples of significant networks include:**

**Social networks**

**World-wide web**

**Internet**

**Air transportation network**

**Cellular network**

**The science of networks is an important, rapidly growing field**

# Networks: basics


---

- A network has  $N$  *nodes* and  $L$  links between nodes
- Each node has a label, e.g. 1, 2, ...,  $N$
- Then a link between node  $i$  and  $j$  can be represented simply as  $(i, j)$

# Networks: basics

---

- A network has  $N$  *nodes* and  $L$  links between nodes
- Each node has a label, e.g. 1, 2, ...,  $N$
- Then a link between node  $i$  and  $j$  can be represented simply as  $(i, j)$




**Example: 6 nodes, 7 links**  
Node one has two edges:  $(1,2)$  and  $(1,5)$

The graph can be represented by the *adjacency matrix*,  $A$   
 $A_{ij}=1$  if there is link between nodes  $i$  and  $j$



# Networks: basics



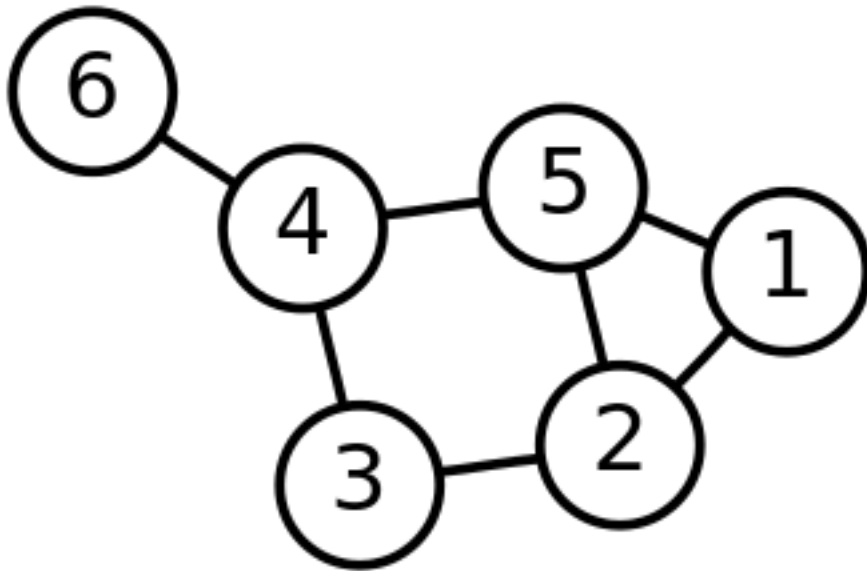
Example: 6 nodes, 7 lines  
Node one has two links: (1,2)  
and (1,5)

The graph can be represented  
by the *adjacency matrix*,  $A$   
 $A_{ij}=1$  if there is link between  
nodes  $i$  and  $j$

$$\begin{pmatrix} 0 & 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \end{pmatrix}$$

$A$  is symmetric.

# Networks: basics



Can also represent connected portions of graph with edge list:

1	1	2	3	4
2	5	3	4	6

Example: 6 nodes, 7 lines  
Node one has two edges: (1,2)  
and (1,5)


The graph can be represented by the *adjacency matrix*,  $A$   
 $A_{ij}=1$  if there is link between nodes  $i$  and  $j$

$$\begin{pmatrix} 0 & 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \end{pmatrix}$$

$A$  is symmetric.

# Networks: basics

---



The *degree* of a node is the number of links connected to it:

$$q_1 = 2, q_5 = 3, \dots$$

The *degree distribution*,  $P(q)$  is particularly important.  $P(q)$  is the fraction of nodes in the graph with degree =  $q$

$$P(1) = 1/6, P(2) = 2/6, P(3) = 3/6$$

**Homework 3:** You will work with a simple model for growing networks.