

## Introduction to High Performance Scientific Computing

Autumn, 2016

Lecture 10

Imperial College  
London

Prasun Ray  
10 November 2016

---

---

---

---

---

---

---

## Today

- A few comments on Fortran
- F2Py with fortran modules
- Timing Fortran code
- Networks

Imperial College  
London

---

---

---

---

---

---

---

## Notes on Fortran

- Do not attempt to guess Fortran syntax! – look it up instead
- Do not develop Fortran code in Python using f2py – get a Fortran-only code working first
- Test code (compile and run) after adding a few (4-5) lines of code. Don't write 100 lines and then test!
- Lecture 7: code structure, variable types, loops, if-then, subroutines
- Lecture 8: allocatable arrays, functions modules
- Lecture 9: Lapack, f2py

Imperial College  
London

---

---

---

---

---

---

---

## Notes on Fortran

### Debugging code

- If the compiler is giving you a series of errors, look at the topmost error message. It will include a line number
- If the code runs but crashes with a segmentation fault: probably a problem with array indices (e.g. trying to access the 12<sup>th</sup> element of a dimension(11) array)
- If you can't tell where the code is crashing, add print statements, e.g. `print *, 1` --some code-- `print *, 2`  
When the code is run, if the 1 prints to screen but not the 2, you know where the problem is.
- If the code runs, but gives the wrong answer, add print statements outputting values of variables, try a small problem size where you know what values the variables can take

Imperial College  
London

## Notes on Fortran

- Modules consist of:
  1. Module variables
  2. Module sub-programs
- Module variables are "available" throughout the module
  - They *do not* need to be declared
  - They *do not* need to be provided as input/output in the sub-program header
- Module variables and module sub-programs are also "available" in any program or sub-program that *uses* the module
- A module by itself doesn't do anything
  - There should be a "main" program which uses it
  - You can compile a module by itself: `gfortran -c module.f90`
  - But to generate an executable, you need a *program*:  
`gfortran -o program.exe module.f90 program.f90`

Imperial College  
London

## Notes on Fortran

- Compiling code that uses lapack routines:

```
$ gfortran -c program.f90
$ gfortran -o program.exe program.o -llapack
or
```

```
$ gfortran -o program.exe program.f90 -llapack
```

Similarly, with f2py:

```
$ f2py -llapack -c program.f90 -m module_name
```

Imperial College  
London

## F2Py and fortran modules

- F2Py will recognize subroutines and functions in modules
- What about variables?
  - Try f2py with circle module from last week (*f2pymodule\_circle.f90*)

```
$ f2py -c f2pymodule_circle.f90 -m cmod
```

```
In [10]: import cmod
```

```
In [11]: cmod.<tab>
cmod.circle  cmod.so
```

Need to look at *cmod.circle*

Imperial College  
London

## F2Py and fortran modules

Need to look at *cmod.circle*:

```
In [12]: cmod.circle?
```

```
Docstring:
'd'-scalar
initialize_pi()
```

```
Wrapper for ``initialize_pi``.
```

```
circumference = circumference(radius)
```

```
Wrapper for ``circumference``.
```

```
Parameters
```

```
radius : input float
```

```
Returns
```

```
circumference : float
area = area(radius)
```

And similar info for "area"

Imperial College  
London

## F2Py and fortran modules

How do we access variables and methods in *cmod.circle*?

```
In [23]: cmod.circle.<tab>
cmod.circle.area  cmod.circle.circumference
cmod.circle.initialize_pi  cmod.circle.pi
```

Can initialize *pi* in python:

```
In [9]: cmod.circle.pi
Out[9]: array(0.0)
```

```
In [10]: cmod.circle.pi = pi
```

```
In [11]: cmod.circle.pi
Out[11]: array(3.141592653589793)
```

Imperial College  
London

## F2Py and fortran modules

How do we access variables and methods in `cmod.circle`?

```
In [23]: cmod.circle.<tab>
cmod.circle.area          cmod.circle.circumference
cmod.circle.initialize_pi  cmod.circle.pi
```

Can initialize `pi` in python:

```
In [9]: cmod.circle.pi
Out[9]: array(0.0)

In [10]: cmod.circle.pi = pi

In [11]: cmod.circle.pi
Out[11]: array(3.141592653589793)
```

Can also initialize allocatable arrays, see `f2pymodule_circle.f90...`

Imperial College  
London

---

---

---

---

---

---

---

---

## F2Py and fortran modules

Can also initialize allocatable arrays, see `f2pymodule_circle_array.f90`:

```
!module for computing circumference, area, and "mass" of circle
module circle
  implicit none
  real(kind=8) :: pi
  real(kind=8), allocatable, dimension(:) :: weights,mass
  save

  ...

  ...

  subroutine compute_mass(radius,mass)
    !compute mass = weights*area
    implicit none
    real(kind=8), intent(in) :: radius
    real(kind=8), intent(out) :: mass(:)

    mass = weights*area(radius)
  end subroutine compute_mass
```

Imperial College  
London

---

---

---

---

---

---

---

---

## F2Py and fortran modules

Basic steps:

1. Compile with f2py: `$ f2py -c f2pymodule_circle_array.f90 -m cmoda`

Imperial College  
London

---

---

---

---

---

---

---

---

## F2Py and fortran modules

### Basic steps:

1. Compile with f2py: `$ f2py -c f2pymodule_circle_array.f90 -m cmoda`

2. Import module in python: `In [4]: import cmoda`

Imperial College  
London

---

---

---

---

---

---

---

---

## F2Py and fortran modules

### Basic steps:

1. Compile with f2py: `$ f2py -c f2pymodule_circle_array.f90 -m cmoda`

2. Import module in python: `In [4]: import cmoda`

3. Initialize variables:

`In [15]: cmoda.circle.pi=pi`

`In [16]: cmoda.circle.weights=arange(5)`

`In [17]: cmoda.circle.pi,cmoda.circle.weights`

`Out[17]: (array(3.141592653589793), array([ 0., 1., 2., 3., 4.]))`

Imperial College  
London

---

---

---

---

---

---

---

---

## F2Py and fortran modules

### Basic steps:

1. Compile with f2py: `$ f2py -c f2pymodule_circle_array.f90 -m cmoda`

2. Import module in python: `In [4]: import cmoda`

3. Initialize variables:

`In [15]: cmoda.circle.pi=pi`

`In [16]: cmoda.circle.weights=arange(5)`

`In [17]: cmoda.circle.pi,cmoda.circle.weights`

`Out[17]: (array(3.141592653589793), array([ 0., 1., 2., 3., 4.]))`

4. Compute mass:

`In [18]: cmoda.circle.compute_mass(2.0)`

`In [19]: cmoda.circle.mass`

`Out[19]: array([ 0. , 12.56637061, 25.13274123, 37.69911184, 50.26548246])`

Imperial College  
London

---

---

---

---

---

---

---

---

### Finite difference methods

- Finite difference methods are a standard approach for numerical differentiation.
- They form the basis for a wide variety of methods used to solve partial differential equations
- See online supplementary class notes: *Notes on numerical differentiation with finite difference methods*

Imperial College  
London

---

---

---

---

---

---

---

---

### Unix *time* command

Can use *time* to obtain timing info for any unix command:

```
$ time ./midpoint.exe
N=      512000
sum=    3.1415926535901515
error=   3.5837999234900053E-013

real    0m0.015s
user    0m0.010s
sys     0m0.003s
```

- real is approximately the wall-clock time
- user is time spent executing the program
- sys is time spent on system tasks required by program

Imperial College  
London

---

---

---

---

---

---

---

---

### Fortran timing functions

- Unix *time* doesn't tell you how much time different parts of program take
- system\_clock and cpu\_time gives wall time and cpu time between two points in code
- See midpoint\_time.f90:

```
!timing variables
real(kind=8) :: cpu_t1,cpu_t2,clock_time
integer(kind=8) :: clock_t1,clock_t2,clock_rate
```

Imperial College  
London

---

---

---

---

---

---

---

---

### Fortran timing functions

- Unix *time* doesn't tell you how much time different parts of program take
- `system_clock` and `cpu_time` gives wall time and cpu time between two points in code
- See `midpoint_time.f90`:  

```
!timing variables
real(kind=8) :: cpu_t1,cpu_t2,clock_time
integer(kind=8) :: clock_t1,clock_t2,clock_rate

call system_clock(clock_t1)
call cpu_time(cpu_t1)
!loop over intervals computing each interval's contribution to
integral
... Midpoint quadrature ...
call cpu_time(cpu_t2)
print *, 'elapsed cpu time (seconds) =',cpu_t2-cpu_t1

call system_clock(clock_t2,clock_rate)
print *, 'elapsed wall time (seconds)= ',
      dble(clock_t2-clock_t1)/dble(clock_rate)
```

Imperial College  
London

### Fortran timing functions

- Unix *time* doesn't tell you how much time different parts of program take
- `system_clock` and `cpu_time` gives wall time and cpu time between two points in code
- See `midpoint_time.f90`:  

```
$ ./midpoint_t.exe
elapsed cpu time (seconds) = 8.623999999999997E-003
elapsed wall time (seconds)= 9.12799966E-03
N= 512000
sum= 3.1415926535901515
error= 3.5837999234900053E-013
```
- Can place timing commands throughout code to find bottlenecks

Imperial College  
London

### Fortran timing functions

- Also, often have a *theoretical estimate* of how cost scales with problem size
- A method may require  $O(N)$  (or  $O(N \ln N)$  or  $O(N^2)$ ) operations
- But does your implementation of the algorithm match theory?
- How do compiler optimizations affect performance?
- Carefully timing code while varying the problem size can help answer these questions

Imperial College  
London

## Networks

Examples of significant networks include:

- Social networks
- World-wide web
- Internet
- Air transportation network
- Cellular network

The science of networks is an important, rapidly growing field

Imperial College  
London

---

---

---

---

---

---

---

---

## Networks: basics

- A network has  $N$  nodes and  $L$  links between nodes
- Each node has a label, e.g. 1, 2, ...,  $N$
- Then a link between node  $i$  and  $j$  can be represented simply as  $(i, j)$

Imperial College  
London

---

---

---

---

---

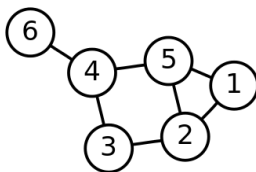
---

---

---

## Networks: basics

- A network has  $N$  nodes and  $L$  links between nodes
- Each node has a label, e.g. 1, 2, ...,  $N$
- Then a link between node  $i$  and  $j$  can be represented simply as  $(i, j)$



Example: 6 nodes, 7 links  
Node one has two edges: (1,2)  
and (1,5)

The graph can be represented  
by the *adjacency matrix*,  $A$   
 $A_{ij}=1$  if there is link between  
nodes  $i$  and  $j$

Imperial College  
London

---

---

---

---

---

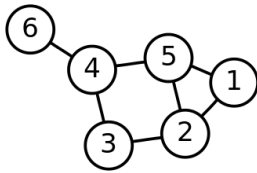
---

---

---



### Networks: basics



Example: 6 nodes, 7 links  
Node one has two links: (1,2)  
and (1,5)

The graph can be represented  
by the *adjacency matrix*,  $A$   
 $A_{ij}=1$  if there is link between  
nodes  $i$  and  $j$

$$\begin{pmatrix} 0 & 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \end{pmatrix}$$

$A$  is symmetric.

Imperial College  
London

---

---

---

---

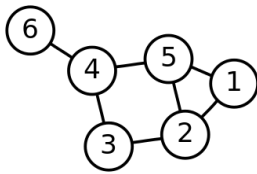
---

---

---

---

### Networks: basics



Example: 6 nodes, 7 links  
Node one has two edges: (1,2)  
and (1,5)

The graph can be represented  
by the *adjacency matrix*,  $A$   
 $A_{ij}=1$  if there is link between  
nodes  $i$  and  $j$

$$\begin{pmatrix} 0 & 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \end{pmatrix}$$

$A$  is symmetric.

Can also represent connected portions  
of graph with edge list:

$$\begin{bmatrix} 1 & 1 & 2 & 2 & 3 & 4 & 4 \\ 2 & 5 & 3 & 5 & 4 & 5 & 6 \end{bmatrix}$$

Imperial College  
London

---

---

---

---

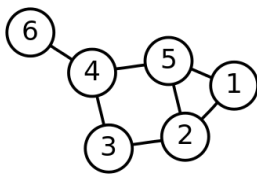
---

---

---

---

### Networks: basics



The *degree* of a node is the  
number of links connected to it:  
 $q_1 = 2, q_5 = 3, \dots$

The *degree distribution*,  $P(q)$  is  
particularly important.  $P(q)$  is  
the fraction of nodes in the  
graph with degree =  $q$

$$P(1) = 1/6, P(2) = 2/6, P(3) = 3/6$$

Homework 3: You will work with a  
simple model for growing networks.

Imperial College  
London

---

---

---

---

---

---

---

---