

# **Introduction to High Performance Scientific Computing**

**Autumn, 2016**

**Lecture 11**

# Today

---

*Further comments on f2py and timing code*

*Compiling and profiling fortran code*

*Introduction to parallel computing*

# Fortran timing functions

---

See `midpoint_time.f90`: timers placed around *do-loop* which computes areas of rectangles

Can use timing results to check if code is behaving “properly.”

```
!timing variables
  real(kind=8) :: cpu_t1,cpu_t2,clock_time
  integer(kind=8) :: clock_t1,clock_t2,clock_rate

  call system_clock(clock_t1)
  call cpu_time(cpu_t1)
  !loop over intervals computing each interval's contribution to
integral
... Midpoint quadrature ...
  call cpu_time(cpu_t2)
  print *, 'elapsed cpu time (seconds) =',cpu_t2-cpu_t1

  call system_clock(clock_t2,clock_rate)
  print *, 'elapsed wall time (seconds)= ',
          dble(clock_t2-clock_t1)/dble(clock_rate)
```

# Fortran timing functions

---

- Run `midpoint_time` with `N=20000` and `N=40000`

```
$ ./midpoint_time.exe
elapsed cpu time (seconds) = 4.44000000000000038E-004
elapsed wall time (seconds)= 5.30000019E-04
N=      20000
```

```
$ ./midpoint_time.exe
elapsed cpu time (seconds) = 8.8399999999999980E-004
elapsed wall time (seconds)= 9.67999978E-04
N=      40000
```

- We can see the cpu time doubles (as we would hope). The wall-time shows more complicated behavior – could be related to what other background processes are running, or just the very short computation time.

# Profiling

---

- Profilers give detailed information about time spent in different parts of code
- In python: `run -p filename` gives profiling info
- With fortran (or c), can use *gprof* utility (not available on Macs)
- Steps:
  1. Compile code with `-pg` flag

```
$ gfortran -pg -o mt2.exe midpoint_time2.f90
```

# Profiling

---

- Profilers give detailed information about time spent in different parts of code
- In python: `run -p filename` gives profiling info
- With fortran (or c), can use *gprof* utility (not available on Macs)
- Steps:
  1. Compile code with `-pg` flag

```
$ gfortran -pg -o mt2.exe midpoint_time2.f90
```

2. Run code (this will generate *gmon.out*):

```
$ ./mt2.exe
```

3. Finally, run *gprof*

```
$ gprof ./mt2.exe
```

# Profiling

---

## Output looks like:

Each sample counts as 0.01 seconds.

%	cumulative	self		self	total	
time	seconds	seconds	calls	s/call	s/call	name
63.67	6.25	6.25	1	6.25	9.85	MAIN__
36.75	9.85	3.61	512000000	0.00	0.00	integrand_

## and:

index	% time	self	children	called	name
		6.25	3.61	1/1	main [2]
[1]	100.0	6.25	3.61	1	MAIN__ [1]
		3.61	0.00	512000000/512000000	integrand_ [3]
-----					
					<spontaneous>
[2]	100.0	0.00	9.85		main [2]
		6.25	3.61	1/1	MAIN__ [1]
-----					
		3.61	0.00	512000000/512000000	MAIN__ [1]
[3]	36.6	3.61	0.00	512000000	integrand_ [3]

-----

# Profiling

---

- Can get line-by-line information from other tools like, *oprof*
- The more complicated the code, the more useful profiling becomes



# Notes on compiling

---

- Up to now:

```
$ gfortran -c program.f90  
$ gfortran -o program.exe program.o -llapack  
or
```

```
$ gfortran -o program.exe program.f90 -llapack
```

- But typically want to turn on *optimization* -O flag:

```
$ gfortran -O3 -c program.f90
```

- -O3 is highest level of optimization (can also use -O1, -O2)

# Notes on compiling

---

- Look at `midpoint_time.f90` compiled with and without `-O3`:

```
$ ./mt2.exe
elapsed cpu time (seconds) = 0.32510499999999998
elapsed wall time (seconds)= 0.326231003
N= 25600
sum= 3.1415926537169634
error= 1.2717027431108363E-010

$ ./mt2_03.exe
elapsed cpu time (seconds) = 0.14376099999999997
elapsed wall time (seconds)= 0.144617006
N= 25600
sum= 3.1415926537169634
error= 1.2717027431108363E-010
```

- Optimization can make a substantial difference
- `f2py` uses `-O3` by default


# Notes on compiling

---

- **Many other useful compiler flags, e.g.:**
  - **-Wall: “warn about all” – generates warnings about common sources of bugs**
  - **-fbounds-check: checks that array index is within bounds of an array (common problem)**
- **Comprehensive list at:**  
**<https://gcc.gnu.org/onlinedocs/gfortran/Invoking-GNU-Fortran.html>**

# Moore's law

## Microprocessor Transistor Counts 1971-2011 & Moore's Law




- Number of transistors on chip doubles every two years
- From ~2005, multicore processors (why?)
- 2 core, then 4, 6, 8, 16
- Intel Xeon Phi: 60+ cores!

# Algorithms and hardware


Speed-Up  
Factor

Derived from Computational Methods



Speed-Up  
Factor

Derived from Supercomputer Hardware



# Why parallelize a code?

---


**1. Serial (single-processor) code is too slow**

**or**

**2 Serial code is too big**

# Parallel computing paradigms

---




## Shared memory

- One 4-core chip with shared memory (RAM)
- MPI can coordinate communication between cores
- OpenMP generally easier to use for shared-memory systems
- MPI = *Message Passing Interface*
- OpenMP = *Open Multi-Processing*

# Parallel computing paradigms

---




## Distributed memory

- Each (4-core) chip has its own memory
- The chips are connected by network 'cables'
- MPI coordinates communication between two or more CPUs



# Parallel computing paradigms

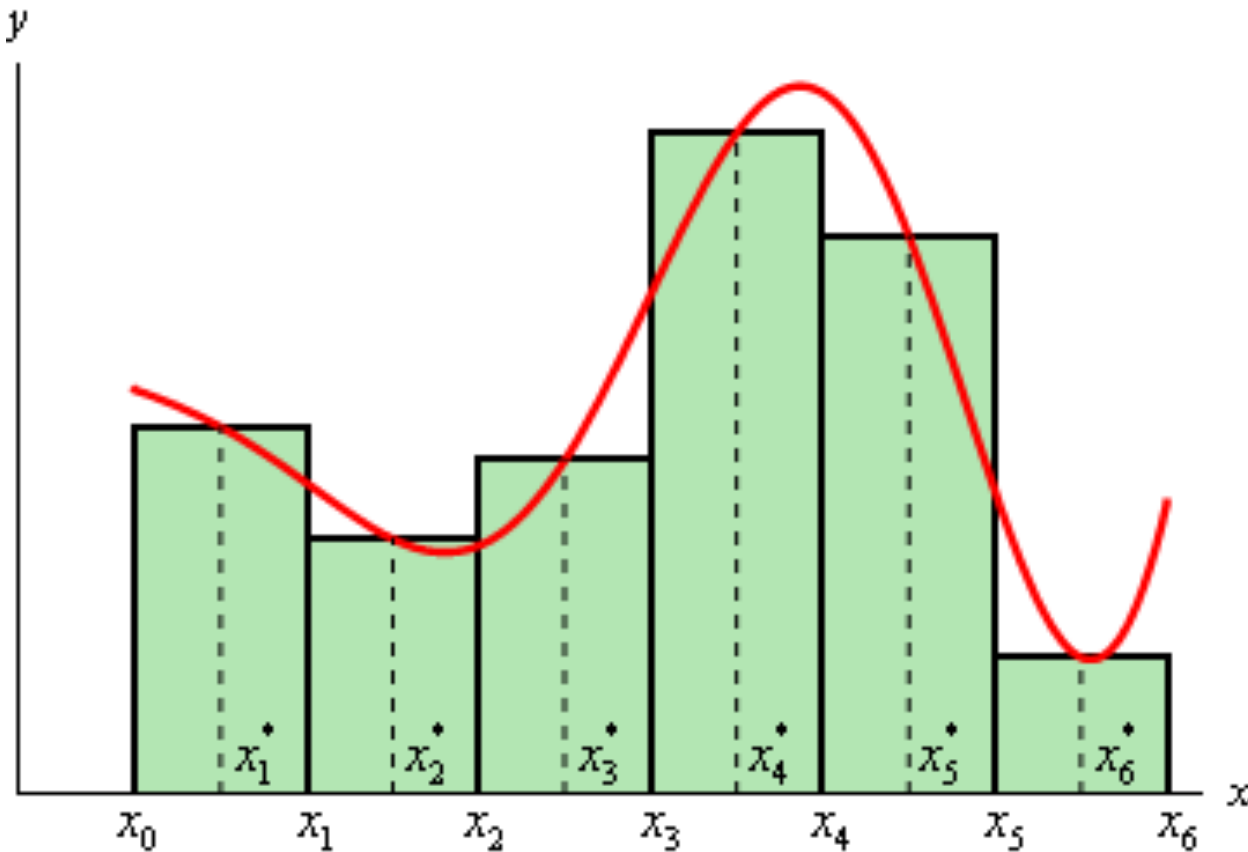
---



## Related approaches:

- Hybrid programming: mix of shared-memory (OpenMP) and distributed-memory (MPI) programming
- GPU's: Shared memory programming (CUDA or OpenCL)
- Coprocessors and co-array programming

# Example: computing an integral



- Estimate integral with midpoint rule,

$$I = \int_{x_0}^{x_6} f(x) dx$$

1. Compute:

$$f(x_1^*), f(x_2^*), \dots$$


2. Compute areas of rectangles:

$$I_1 = (x_1 - x_0) * f(x_1^*)$$

3. Sum areas:

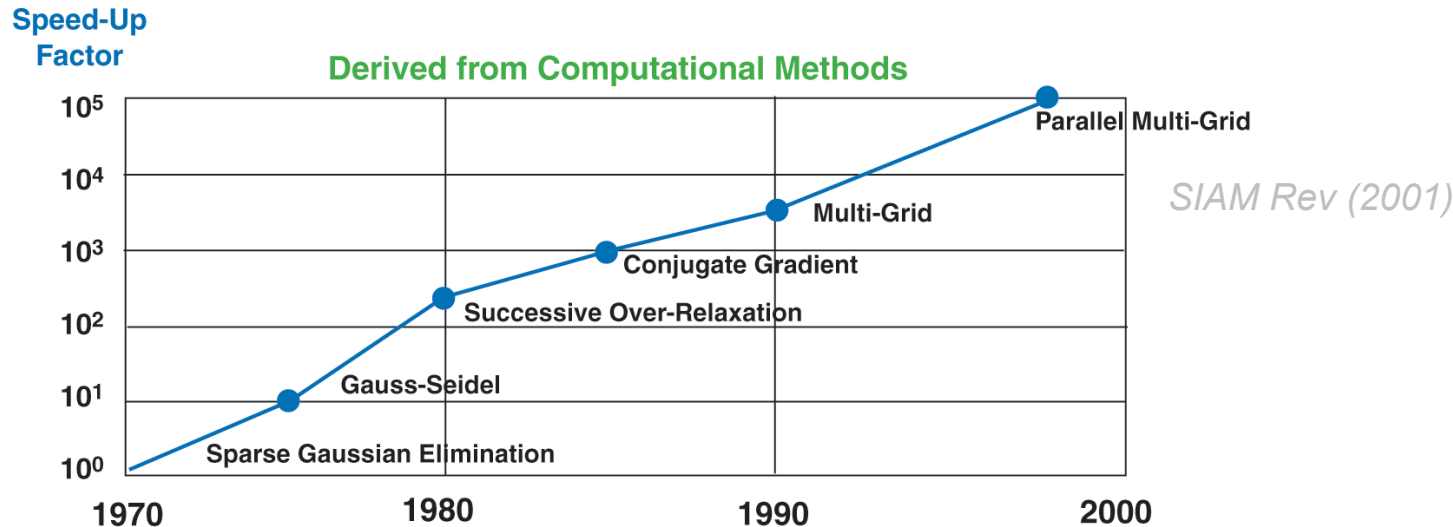
$$I \approx I_1 + I_2 + I_3 + \dots$$

# Example: computing an integral




- How to parallelize?
- With three processors, can compute areas of two rectangles on each processor
- Not practical for small calculations, but could split  $1e7$  rectangles across, say, 10 processors

# Scaling and performance



- How do we measure performance of a parallel code?
- Serial code: Optimize the *efficiency* → cost required to obtain a certain level of accuracy

# Scaling and performance



- How do we measure performance of a parallel code?
- Serial code: Optimize the *efficiency* → cost required to obtain a certain level of accuracy
- Parallel code: Also optimize *scaling* or *speedup*: how much faster is the calculation when the number of procs is increased?

# Speedup

---

- **Speedup = Computation time on one proc/time on N procs =  $T_s/T_p$**
- **Ideal: N = 10 processors, speedup = N = 10**

# Speedup

---

- **Speedup = Computation time on one proc/time on N procs =  $T_s/T_p$**
- **Ideal:  $N = 10$  processors, speedup =  $N = 10$**
- **Real life: Speedup will be less than  $N$  (possibly much less) Why?**
  - **Startup costs**
  - **Communication**
  - **Only part of the algorithm parallelizes**
- **Typically interested in performance of large problems running on large number of processors**
  - **Workstation:  $N = 16, 32$**
  - **Imperial HPC (cx2):  $N = 256+$**
  - **UK HPC (Archer):  $N = 1e3, 1e4, \dots$**
- **Ahmdal's law provides guidance**

# Ahmdal's law

---

- Usually only part of a computation can be parallelized
  - One processor:  $T(1) = s + p$
  - Two processors:  $T(2) = s + p/2$
  - N processors:  $T(N) = s + p/N$

p is the part of the code that can be parallelized



# Ahmdal's law

---

- Usually only part of a computation can be parallelized
  - One processor:  $T(1) = s + p$
  - Two processors:  $T(2) = s + p/2$
  - N processors:  $T(N) = s + p/N$

$p$  is the part of the code that can be parallelized

So, if only half the code can be parallelized ( $s = p = 0.5$ ), Then the maximum speedup  $T(1)/T(N \rightarrow \infty) = (s+p)/(s) = 2$

It is important for  $p$  to be much larger than  $s$ !

# Ahmdal's law

---

Speedup  $T(1)/T(N) = (s+p)/(s+p/N)$

Example:  $s = 0.1$ ,  $p = 0.9$

Number of processors	Speedup
1	1
2	1.8
4	3.1
8	4.7
16	6.4
32	7.8
256	9.7

# Ahmdal's law

---

**Speedup  $T(1)/T(N) = (s+p)/(s+p/N)$**

**Example:  $s = 0.1$ ,  $p = 0.9$**

Number of processors	Speedup
1	1
2	1.8
4	3.1
8	4.7
16	6.4
32	7.8
256	9.7

**Waste of resources to use  $N=256$ !**

# Strong and weak scaling

---

- **Strong scaling:** Time needed to solve a problem of fixed size as number of processors increases
- **Weak scaling:** Time needed for problem with *fixed size per processor*