

Introduction to High Performance Scientific Computing

Autumn, 2016

Lecture 11

Imperial College
London

Prasun Ray
14 November 2016

Today

Further comments on f2py and timing code

Compiling and profiling fortran code

Introduction to parallel computing

Imperial College
London

Fortran timing functions

See `midpoint_time.f90`: timers placed around *do-loop* which computes areas of rectangles

Can use timing results to check if code is behaving "properly."

```
!timing variables
real(kind=8) :: cpu_t1,cpu_t2,clock_time
integer(kind=8) :: clock_t1,clock_t2,clock_rate

call system_clock(clock_t1)
call cpu_time(cpu_t1)
!loop over intervals computing each interval's contribution to
integral
... Midpoint quadrature ...
call cpu_time(cpu_t2)
print *, 'elapsed cpu time (seconds) =',cpu_t2-cpu_t1

call system_clock(clock_t2,clock_rate)
print *, 'elapsed wall time (seconds)= ',
      dble(clock_t2-clock_t1)/dble(clock_rate)
```

Imperial College
London

Fortran timing functions

- Run `midpoint_time` with `N=20000` and `N=40000`

```
$ ./midpoint_time.exe
elapsed cpu time (seconds) = 4.4400000000000038E-004
elapsed wall time (seconds)= 5.30000019E-04
N= 20000
```

```
$ ./midpoint_time.exe
elapsed cpu time (seconds) = 8.8399999999999980E-004
elapsed wall time (seconds)= 9.67999978E-04
N= 40000
```

- We can see the cpu time doubles (as we would hope). The wall-time shows more complicated behavior – could be related to what other background processes are running, or just the very short computation time.

Imperial College
London

Profiling

- Profilers give detailed information about time spent in different parts of code
- In python: `run -p filename` gives profiling info
- With fortran (or c), can use *gprof* utility (not available on Macs)
- Steps:
 1. Compile code with `-pg` flag


```
$ gfortran -pg -o mt2.exe midpoint_time2.f90
```

Imperial College
London

Profiling

- Profilers give detailed information about time spent in different parts of code
- In python: `run -p filename` gives profiling info
- With fortran (or c), can use *gprof* utility (not available on Macs)
- Steps:
 1. Compile code with `-pg` flag


```
$ gfortran -pg -o mt2.exe midpoint_time2.f90
```
 2. Run code (this will generate *gmon.out*):


```
$ ./mt2.exe
```
 3. Finally, run *gprof*

```
$ gprof ./mt2.exe
```

Imperial College
London

Profiling

Output looks like:

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self s/call	total s/call	name
63.67	6.25	6.25	1	6.25	9.85	MAIN__
36.75	9.85	3.61	512000000	0.00	0.00	integrand_

and:

index	% time	self	children	called	name
		6.25	3.61	1/1	main [2]
[1]	100.0	6.25	3.61	1	MAIN__ [1]
		3.61	0.00	512000000/512000000	integrand_ [3]

[2]	100.0	0.00	9.85		<spontaneous>
		6.25	3.61	1/1	main [2]
					MAIN__ [1]

[3]	36.6	3.61	0.00	512000000/512000000	MAIN__ [1]
		3.61	0.00	512000000	integrand_ [3]

Imperial College
London

Profiling

- Can get line-by-line information from other tools like, *oprof*
- The more complicated the code, the more useful profiling becomes

Imperial College
London

Notes on compiling

- Up to now:

```
$ gfortran -c program.f90
$ gfortran -o program.exe program.o -llapack
or
```

```
$ gfortran -o program.exe program.f90 -llapack
```

- But typically want to turn on *optimization* -O flag:

```
$ gfortran -O3 -c program.f90
```

- -O3 is highest level of optimization (can also use -O1, -O2)

Imperial College
London

Notes on compiling

- Look at midpoint_time.f90 compiled with and without -O3:

```
$ ./mt2.exe
elapsed cpu time (seconds) = 0.32510499999999998
elapsed wall time (seconds)= 0.326231003
N= 25600
sum= 3.1415926537169634
error= 1.2717027431108363E-010
```

```
$ ./mt2_O3.exe
elapsed cpu time (seconds) = 0.14376099999999997
elapsed wall time (seconds)= 0.144617006
N= 25600
sum= 3.1415926537169634
error= 1.2717027431108363E-010
```

- Optimization can make a substantial difference
- f2py uses -O3 by default

Imperial College
London

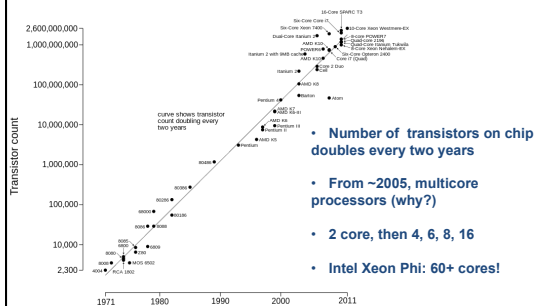
Notes on compiling

- Many other useful compiler flags, e.g.:
 - Wall: "warn about all" – generates warnings about common sources of bugs
 - fbounds-check: checks that array index is within bounds of an array (common problem)
- Comprehensive list at:
<https://gcc.gnu.org/onlinedocs/gfortran/Invoking-GNU-Fortran.html>

Imperial College
London

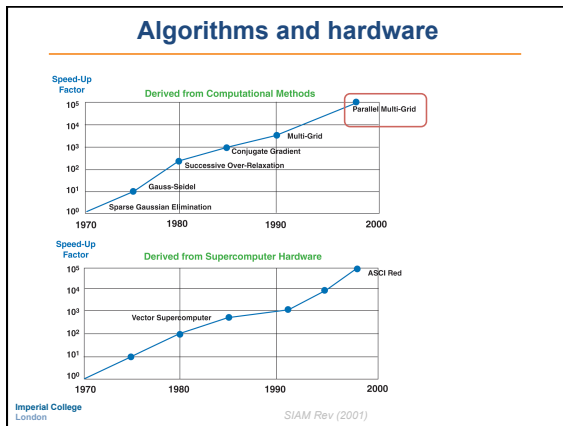
Moore's law

Microprocessor Transistor Counts 1971-2011 & Moore's Law



Imperial College
London

"Transistor Count and Moore's Law - 2011" by Wgislmon - Own work. Licensed under CC BY-SA 3.0 via Wikimedia Commons - http://commons.wikimedia.org/wiki/File:Transistor_Count_and_Moore%27s_Law_-_2011.svg#/media/File:Transistor_Count_and_Moore%27s_Law_-_2011.svg



Why parallelize a code?

1. Serial (single-processor) code is too slow
- or
2. Serial code is too big

Imperial College London

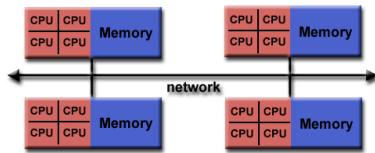
Parallel computing paradigms

Shared memory

- One 4-core chip with shared memory (RAM)
- MPI can coordinate communication between cores
- OpenMP generally easier to use for shared-memory systems
- MPI = *Message Passing Interface*
- OpenMP = *Open Multi-Processing*

Imperial College London

Parallel computing paradigms

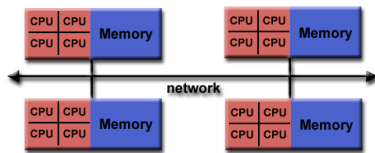


Distributed memory

- Each (4-core) chip has its own memory
- The chips are connected by network 'cables'
- MPI coordinates communication between two or more CPUs

Imperial College
London

Parallel computing paradigms

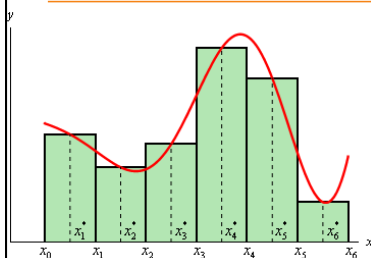


Related approaches:

- Hybrid programming: mix of shared-memory (OpenMP) and distributed-memory (MPI) programming
- GPU's: Shared memory programming (CUDA or OpenCL)
- Coprocessors and co-array programming

Imperial College
London

Example: computing an integral



- Estimate integral with midpoint rule,

$$I = \int_{x_0}^{x_6} f(x) dx$$

1. Compute:

$$f(x_1^*), f(x_2^*), \dots$$

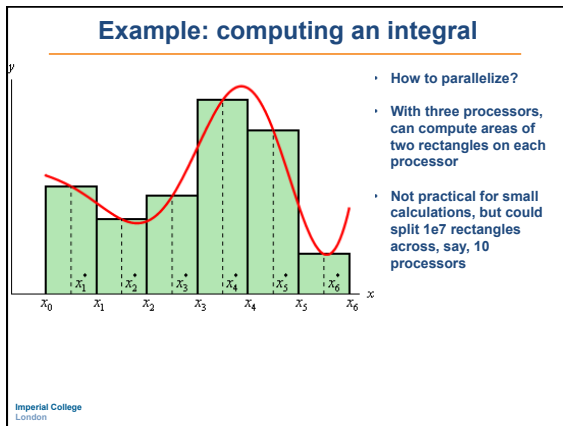
2. Compute areas of rectangles:

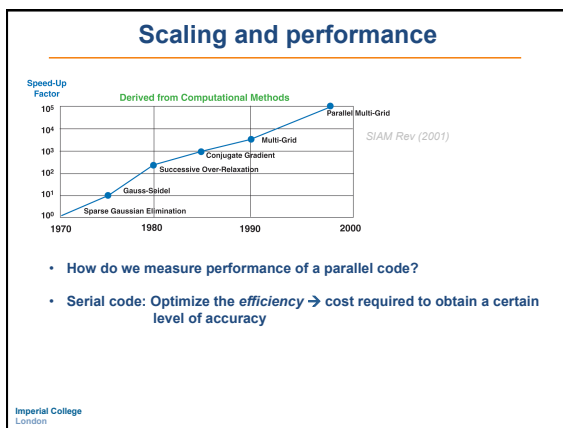
$$I_1 = (x_1 - x_0) * f(x_1^*)$$

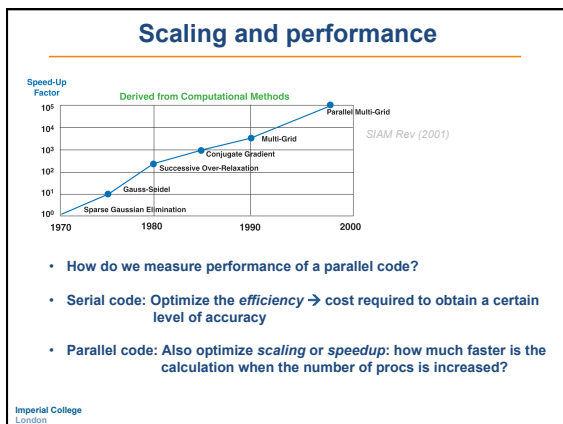
3. Sum areas:

$$I \approx I_1 + I_2 + I_3 + \dots$$

Imperial College
London







Speedup

- Speedup = Computation time on one proc/time on N procs = T_s/T_p
- Ideal: N = 10 processors, speedup = N = 10

Imperial College
London

Speedup

- Speedup = Computation time on one proc/time on N procs = T_s/T_p
- Ideal: N = 10 processors, speedup = N = 10
- Real life: Speedup will be less than N (possibly much less) Why?
 - Startup costs
 - Communication
 - Only part of the algorithm parallelizes
- Typically interested in performance of large problems running on large number of processors
 - Workstation: N= 16, 32
 - Imperial HPC (cx2): N = 256+
 - UK HPC (Archer): N = 1e3, 1e4, ...
- Ahmdal's law provides guidance

Imperial College
London

Ahmdal's law

- Usually only part of a computation can be parallelized
 - One processor: $T(1) = s + p$
 - Two processors: $T(2) = s + p/2$
 - N processors: $T(N) = s + p/N$
- p is the part of the code that can be parallelized

Imperial College
London

Ahmdal's law

- Usually only part of a computation can be parallelized

- One processor: $T(1) = s + p$
- Two processors: $T(2) = s + p/2$
- N processors: $T(N) = s + p/N$

p is the part of the code that can be parallelized

So, if only half the code can be parallelized ($s = p = 0.5$), Then the maximum speedup $T(1)/T(N \rightarrow \infty) = (s+p)/(s) = 2$

It is important for p to be much larger than s !

Imperial College
London

Ahmdal's law

Speedup $T(1)/T(N) = (s+p)/(s+p/N)$

Example: $s = 0.1$, $p = 0.9$

Number of processors	Speedup
1	1
2	1.8
4	3.1
8	4.7
16	6.4
32	7.8
256	9.7

Imperial College
London

Ahmdal's law

Speedup $T(1)/T(N) = (s+p)/(s+p/N)$

Example: $s = 0.1$, $p = 0.9$

Number of processors	Speedup
1	1
2	1.8
4	3.1
8	4.7
16	6.4
32	7.8
256	9.7

Waste of resources to use $N=256$!

Imperial College
London

Strong and weak scaling

- Strong scaling: Time needed to solve a problem of fixed size as number of processors increases
- Weak scaling: Time needed for problem with *fixed size per processor*

Imperial College
London
