

# **Introduction to High Performance Scientific Computing**

**Autumn, 2016**

**Lecture 12**

# Today

---

*Introduction to OpenMP*

*Getting started*

*Parallel regions*

*Parallel loops*

# Overview

---

- OpenMP provides a fairly easy approach to parallelizing c/c++ or fortran code
- Add *directives* indicating how/where the code should run in parallel
- Parallel regions have multiple threads, each of which should be assigned computational tasks
  - OpenMP is for *shared-memory* parallel programming
  - Each thread has access to all variables that existed before parallel region was created
  - This can cause problems if multiple threads try to change the same variable!
- Particularly useful for parallelizing loops
- When compiling, add *-fopenmp* flag

# Overview

---

Program starts with  
single *master thread*

Start program



master thread


# Overview

---

Program starts with  
single *master thread*

Then, launch parallel  
region with multiple  
threads.

Each thread has  
access to all  
variables introduced  
previously



# Overview

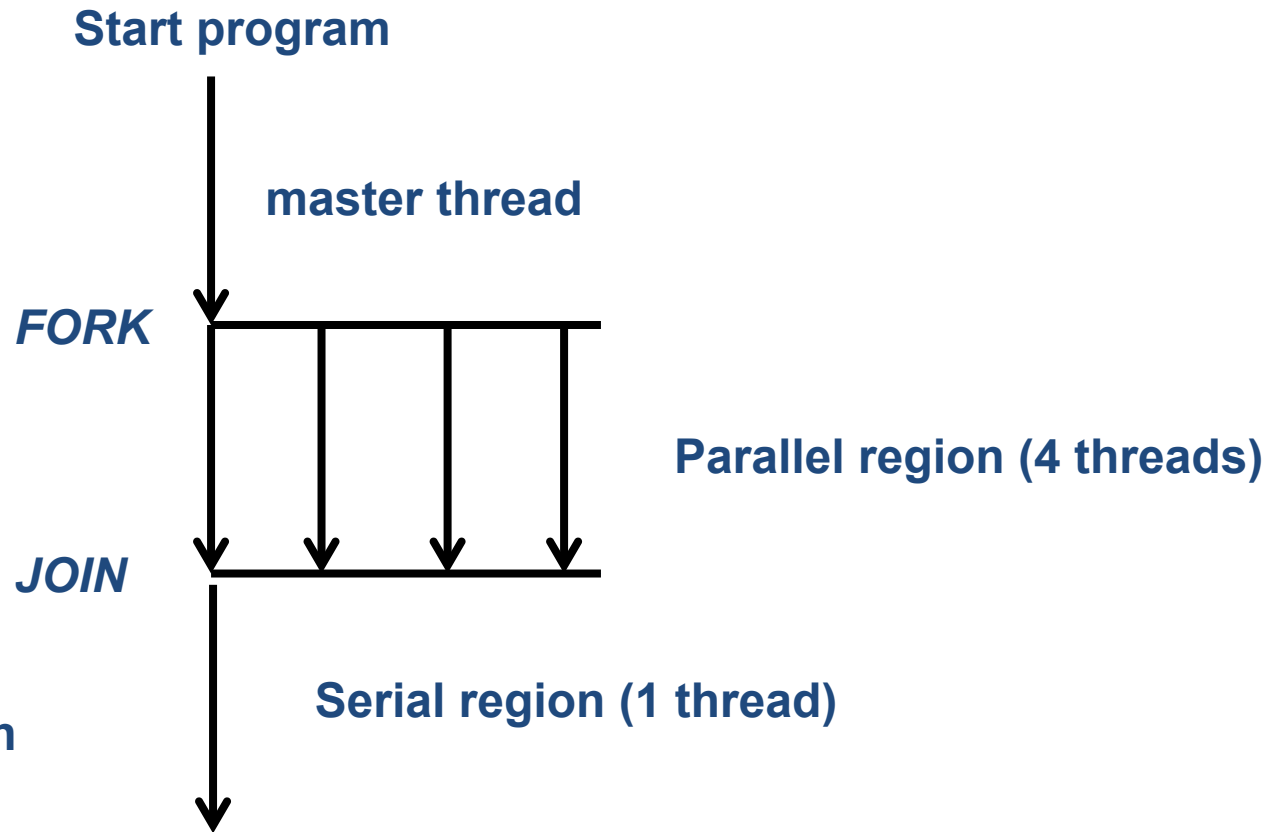
---

Program starts with single *master thread*

Then, launch parallel region with multiple threads.

Each thread has access to all variables introduced previously

Can end parallel region if/when desired and launch parallel regions again in future as needed



# Simple OpenMP example

---

- Launch parallel region, get info on threads (see *firstomp\_v0.f90*)
- Must use openMP module, *omp\_lib*
  - This makes functions like *omp\_get\_num\_threads* available

!Getting started with OpenMP

```
program firstomp
  use omp_lib !makes OpenMP routines, variables available
  implicit none
  integer :: NumThreads,threadID
```

# Simple OpenMP example

---

- Launch parallel region, get info on threads (see *firstomp\_v0.f90*)
- Must use openMP module, *omp\_lib*
  - This makes functions like *omp\_get\_num\_threads* available

!Getting started with OpenMP

```
program firstomp
  use omp_lib !makes OpenMP routines, variables available
  implicit none
  integer :: NumThreads, threadID
!$OMP PARALLEL
  NumThreads = omp_get_num_threads()
  threadID = omp_get_thread_num()
  print *, 'this is thread', threadID, ' of ', NumThreads
!$OMP END PARALLEL
```

- !\$OMP starts an OpenMP directive (#pragma omp in c)
- !\$OMP PARALLEL starts a parallel region (forks a number of threads)
  - *omp\_get\_num\_threads* tells us how many threads are forked
  - *omp\_get\_thread\_num* tells us which thread is being used



# Simple OpenMP example

---

- **Let's compile and run this:**

```
$ gfortran -fopenmp -o testv0.exe firstomp_v0.f90
```

```
$ ./testv0.exe
```

```
this is thread      1  of      4  
this is thread      1  of      4  
this is thread      1  of      4  
this is thread      1  of      4
```

- **Total number of threads is correct, but problem getting the thread id.**

# Simple OpenMP example

---

- **Let's compile and run this:**

```
$ gfortran -fopenmp -o testv0.exe firstomp_v0.f90
```

```
$ ./testv0.exe
```

```
this is thread      1  of      4  
this is thread      1  of      4  
this is thread      1  of      4  
this is thread      1  of      4
```

- **Total number of threads is correct, but problem getting the thread id.**
- **Remember: *threadID* is a *shared* variable.**
  - **Each thread is writing to the same variable with it's id, so only the “last” thread has its ID displayed**

# Simple OpenMP example

---

- **Let's compile and run this:**

```
$ gfortran -fopenmp -o testv0.exe firstomp_v0.f90
```

```
$ ./testv0.exe
```

```
this is thread      1  of      4
this is thread      1  of      4
this is thread      1  of      4
this is thread      1  of      4
```

- **Total number of threads is correct, but problem getting the thread id.**
- **Remember: *threadID* is a *shared* variable.**
  - **Each thread is writing to the same variable with it's id, so only the “last” thread has its ID displayed**
- **How can we fix this? First approach: define a *critical* region...**

# Simple OpenMP example

---

- A critical region, defined with `!$OMP CRITICAL`, runs in serial
  - The threads carry out their tasks sequentially (*firstomp\_v1.f90*)

```
!$OMP PARALLEL
  NumThreads = omp_get_num_threads()
  !$OMP CRITICAL
    threadID = omp_get_thread_num()
    print *, 'this is thread', threadID, ' of ', NumThreads
  !$OMP END CRITICAL
!$OMP END PARALLEL
```

# Simple OpenMP example

- A critical region, defined with `!$OMP CRITICAL`, runs in serial
  - The threads carry out their tasks sequentially (*firstomp\_v1.f90*)

```
!$OMP PARALLEL
  NumThreads = omp_get_num_threads()
  !$OMP CRITICAL
    threadID = omp_get_thread_num()
    print *, 'this is thread', threadID, ' of ', NumThreads
  !$OMP END CRITICAL
!$OMP END PARALLEL
```

So now, if we compile and run:

```
$ ./testv1.exe
this is thread          0  of          4
this is thread          2  of          4
this is thread          1  of          4
this is thread          3  of          4
```

# Simple OpenMP example

---

- Critical regions useful for: data I/O, displaying results
- But forcing serial execution in a parallel region, not generally desirable
- Better approach: set `threadID` to be a *private* variable (*firstomp.f90*)
  - Then, each thread will have their own private copy of the variable

# Simple OpenMP example

---

- Critical regions useful for: data I/O, displaying results
- But forcing serial execution in a parallel region, not generally desirable
- Better approach: set threadID to be a *private* variable (*firstomp.f90*)
  - Then, each thread will have their own private copy of the variable

```
!$OMP PARALLEL PRIVATE(threadID)
  NumThreads = omp_get_num_threads()
  threadID = omp_get_thread_num()
  print *, 'this is thread', threadID, ' of ', NumThreads
!$OMP END PARALLEL
```

```
$ ./test.exe
this is thread      0  of      4
this is thread      2  of      4
this is thread      1  of      4
this is thread      3  of      4
```

# Simple parallel calculation

---

- Can use *threadID* to assign tasks to threads:

```
!$OMP PARALLEL PRIVATE(threadID)
  NumThreads = omp_get_num_threads()
  threadID = omp_get_thread_num()
```

```
  if (threadID==0) then
    call subroutine1(in1,out1)
  elseif (threadID==1) then
    call subroutine1(in2,out2)
  end if
```

```
!$OMP END PARALLEL
```

- Important to distribute work evenly across threads (load balancing)



# Overview

---

- **OpenMP (primarily) consists of *directives* and *routines***
- **Directives are denoted with !OMP**
  - !OMP parallel, !OMP critical, ...
  - **Directives are recognized when `–fopenmp compile-flag` is used**
  - **Otherwise, they are interpreted as comments**
    - **What happens if you use:**  
!\$ print \*, “compiled with `–fopenmp`”
- **Routines are available via the `use omp_lib` command**
  - **e.g. `omp_get_thread_num` and `omp_get_num_threads`**

# Parallel loops

---

- Loops form the backbone of most scientific codes
- They should be parallelized whenever possible
- They can be parallelized if the calculations of each iterations are independent of each other (no data dependencies)

```
do i1 = 1,n  
    x(i1) = y(i1) + z(i1)  
end do
```

Ok to parallelize

```
do i1 = 1,n  
    norm = norm + abs(x(i1))  
end do
```

Can't parallelize easily: each thread updating, *norm*

# Parallel loops

---

- OpenMP makes it very easy to parallelize loops

```
!$OMP parallel do
do i1 = 1,n
    x(i1) = y(i1) + z(i1)
end do
!$OMP end parallel do
```

- OpenMP automatically distributes iterations across threads
  - If NumThreads=2 and n=10, iterations 1,...,5 would be given to thread 0 and iterations 6,...,10 would be done by thread 1 (or vice versa)
  - The iterated variable, *i1*, is automatically set to *private*. Each thread has its own copy.

# Parallel loops

---

- Simple example (*loop\_omp1.f90*):

```
!$OMP parallel do private(threadID)
do i1 = 1,size(x)
    x(i1) = y(i1) + z(i1)
    threadID = omp_get_thread_num()
    print *, 'iteration ',i1,' assigned to thread ',threadID
end do
!$OMP end parallel do

print *, 'test:', maxval(abs(x-y-z))
```

- Note: *threadID* again set to *private*
- Compile and run...

# Parallel loops

---

- Simple example (*loop\_omp1.f90*):

```
!$OMP parallel do private(threadID)
do i1 = 1,size(x)
    x(i1) = y(i1) + z(i1)
    threadID = omp_get_thread_num()
    print *, 'iteration ',i1,' assigned to thread ',threadID
end do
!$OMP end parallel do
```

```
print *, 'test:', maxval(abs(x-y-z))
```

```
$ gfortran -fopenmp -o testl1.exe loop_omp1.f90
```

```
$ ./testl1.exe
```

iteration	1	assigned to thread	0
iteration	3	assigned to thread	2
iteration	2	assigned to thread	1
iteration	4	assigned to thread	3
test:	2.2204460492503131E-016		

# Parallel loops

---

- Can easily “embed” parallel loop in parallel region:

```
!$OMP parallel
!$OMP do
do i1 = 1,n
    x(i1) = y(i1) + z(i1)
end do
!$OMP end do

!Other parallel calculations

!$OMP end parallel
```

# Parallel loops

---

- Now, return to “unparallelizable” example:

```
do i1 = 1,n  
    norm = norm + abs(x(i1))  
end do
```

- Next lecture: use *reduction* to parallelize
- Now:
  1. Let each thread have it's own copy of norm
  2. Sum each thread's partial sum in *critical* region

# Parallel loops

---

## 1. Let each thread have it's own copy of norm

```
norm=0.d0
partial_norm=0.d0
!$OMP parallel firstprivate(partial_norm),private(threadID)
!$OMP do
do i1 = 1,size(x)
    partial_norm = partial_norm + abs(x(i1))
end do
!$OMP end do
```

- The partial sum, *partial\_norm*, is a private variable which must be initialized
- *firstprivate* initializes each thread's value to the value set before the parallel region



# Parallel loops

---

1. Let each thread have it's own copy of norm
2. Sum each thread's partial sum in *critical* region (see *norm\_omp1.f90*)

```
norm=0.d0
partial_norm=0.d0
!$OMP parallel firstprivate(partial_norm),private(threadID)
!$OMP do
do i1 = 1,size(x)
    partial_norm = partial_norm + abs(x(i1))
end do
!$OMP end do

!$OMP critical
threadID = omp_get_thread_num()
print *, 'Thread number:',threadID, 'partial norm=',partial_norm
norm = norm + partial_norm
!$OMP end critical
```

# Parallel loops

---

Compile and run, testing code with:

```
print *, 'test:', norm-sum(abs(x))
```

```
$ gfortran -fopenmp -o testn1.exe norm_omp1.f90
```

```
$ ./testn1.exe
```

```
Thread number:      2 partial norm=    3.1411200080598674
Thread number:      3 partial norm=    3.2431975046920716
Thread number:      1 partial norm=    2.9092974268256819
Thread number:      0 partial norm=    1.8414709848078965
test:    0.000000000000000000
```

# Parallel loops: nested loops

---

- Often work with nested loops:

```
do j1 = 1,N
  do i1 = 1,M
    x(i1,j1) = y(i1,j1) + z(i1,j1)
  end do
end do
```

- Should we parallelize the inner or outer loop? (assuming  $M \sim N$ )

```
!$OMP parallel do private(i1)
do j1 = 1,N
  do i1 = 1,M
    x(i1,j1) = y(i1,j1) + z(i1,j1)
  end do
end do
!$OMP end parallel do
```

outer

# Parallel loops: nested loops

---

- Often work with nested loops:

```
do j1 = 1,N
  do i1 = 1,M
    x(i1,j1) = y(i1,j1) + z(i1,j1)
  end do
end do
```

- Should we parallelize the inner or outer loop? (assuming  $M \sim N$ )

```
do j1 = 1,N
  !$OMP parallel do private
    do i1 = 1,M
      x(i1,j1) = y(i1,j1) + z(i1,j1)
    end do
  !$OMP end parallel do
end do
```

inner

# Parallel loops: nested loops

---

- Should we parallelize the inner or outer loop? (assuming  $M \sim N$ )

```
!$OMP parallel do private(i1)
do j1 = 1,N
  do i1 = 1,M
    x(i1,j1) = y(i1,j1) + z(i1,j1)
  end do
end do
!$OMP end parallel do
```

**Better to parallelize outer thread (setting the inner variable,  $i1$ , to private)**

**If inner loop is parallelized: forking/joining of threads is repeated with each outer loop: this is inefficient!**

# Parallel loops: nested loops

---

Must always be sure loop(s) can be parallelized

Example:

```
!$OMP parallel do private(i1)
do j1 = 2,N
    do i1 = 1,M
        x(i1,j1) = x(i1,j1-1)
    end do
end do
!$OMP end parallel do
```

**Incorrect**

- Different  $j1$ 's are assigned to different threads
- $x(i1,j1-1)$  may not have been computed at the time that it is needed

# Parallel loops: nested loops

---

Must always be sure loop(s) can be parallelized

Example:

```
!$OMP parallel do private(j1)
do i1 = 1,M
    do j1 = 2,N
        x(i1,j1) = x(i1,j1-1)
    end do
end do
!$OMP end parallel do
```

Correct

- Solution: swap inner and outer loops
- Now, computation of  $x$  is “safe.”
  - The “ $i1$  loop” is parallelized, and calculations of  $x$  do not depend on the order in which  $i1$  is iterated.