

Introduction to High Performance Scientific Computing

Autumn, 2016

Lecture 13

Imperial College
London

Prasun Ray
21 November 2016

Today

More on OpenMP

Reductions

Setting number of threads

A few useful OpenMP commands

Imperial College
London

Last time: Parallel loops

Must always be sure loop(s) can be parallelized

Example:

```
!$OMP parallel do private(i1)
do j1 = 2,N
  do i1 = 1,M
    x(i1,j1) = x(i1,j1-1)
  end do
end do
!$OMP end parallel do
```

Is the order of the iterations important? (data dependency)

Do different iterations assign values to same variable? (race condition)

Imperial College
London

Last time: Parallel loops

Must always be sure loop(s) can be parallelized

Example:

```
!$OMP parallel do private(j1)
do i1 = 1,M
  do j1 = 2,N
    x(i1,j1) = x(i1,j1-1)
  end do
end do
!$OMP end parallel do
```

Imperial College
London

Example: computing a norm

Last time: developed simple code for computing norm: $\sum(|x|)$

Serial version:

```
do i1 = 1,size(x)
  norm = norm + abs(x(i1))
end do
```

Imperial College
London

Example: computing a norm

Last time: developed simple code for computing norm: $\sum(|x|)$

Serial version:

```
do i1 = 1,size(x)
  norm = norm + abs(x(i1))
end do
```

Parallel version:

```
!$OMP parallel firstprivate(partial_norm)
!$OMP do
do i1 = 1,size(x)
  partial_norm = partial_norm + abs(x(i1))
end do
!$OMP end do

!$OMP critical
norm = norm + partial_norm
!$OMP end critical
!$OMP end parallel
```

Example: computing a norm

- Typically want to avoid using critical regions
- reduction* provides a simpler approach:

```
!$OMP parallel do reduction(+:norm)
do i1 = 1,size(x)
  norm = norm + abs(x(i1))
end do
!$OMP end parallel do
```

Imperial College
London

Example: computing a norm

- Typically want to avoid using critical regions
- reduction* provides a simpler approach (*omp_norm2.f90*):

```
!$OMP parallel do reduction(+:norm)
do i1 = 1,size(x)
  norm = norm + abs(x(i1))
end do
!$OMP end parallel do
```

- Generally, *reduction* "reduces" an array of numbers distributed across multiple threads to a single number
- Several operations are available, a few common operators are: +, -, *, max, min, and, or.
- Not specific to OpenMP! In MPI, we will use *MPI_REDUCE*.
- Due to ease-of-use and usefulness, one of the most important tools in parallel computing!

Imperial College
London

Example: reduction with *min*

- Here, computation of *x* is parallelized
- Reduction* is used to find $\min(|x|)$

```
!$OMP parallel do reduction(min:xmin)
do i1=1,size(x)
  x(i1) = z(i1)+y(i1)
  xmin = min(abs(x(i1)),xmin)
end do
!$OMP end parallel do
```

Imperial College
London

Setting number of threads

- By default, OpenMP "detects" the number of threads on computer and uses all of them
- Can also set threads in two ways:
 - Within code with `omp_set_num_threads`, e.g.:

```
!$ call omp_set_num_threads(2)
```

(the "\$" ensures this is only called if `-fopenmp` flag is used when compiling)
 - From Unix terminal before program execution:

```
$ export OMP_NUM_THREADS=2
```

Imperial College
London

Other useful OpenMP directives

- Consider a parallel region of code:


```
!$OMP parallel
!code run by *each* thread
!$OMP end parallel
```
- There are a number of directives which we can use in the parallel region

Imperial College
London

Other useful OpenMP directives

- do-loops


```
!$OMP parallel private(i1)
do i1=1,N
!some operations
end do
!$OMP end parallel
```
- In the example above, the full do-loop is run by each thread

Imperial College
London

Other useful OpenMP directives

- do-loops

```
!$OMP parallel private(i1)
do i1=1,N
  !some operations
end do
!$OMP end parallel
```

- In the example above, the full do-loop is run by each thread

```
!$OMP parallel private(i1)
!$OMP do
do i1=1,N
  !some operations
end do
!$OMP end do
!$OMP end parallel
```

- Now, this is the same as a parallel do loop

Imperial College
London

Other useful OpenMP directives

- Sections

```
!$OMP parallel
!$OMP sections
!$OMP section
  !code run by one thread
!$OMP section
  !code run by second thread
!$OMP section
  !code run by another thread
!$OMP end sections
!$OMP end parallel
```

- Manually assign tasks to threads
- For example, invert four matrices (of the same size)
- Could have four "sections", one for each matrix inversion

Imperial College
London

Last lecture: Simple parallel calculation

- Can use *threadID* to assign tasks to threads:

```
!$OMP PARALLEL PRIVATE(threadID)
NumThreads = omp_get_num_threads()
threadID = omp_get_thread_num()
```

```
if (threadID==0) then
  call subroutine1(in1,out1)
elseif (threadID==1) then
  call subroutine1(in2,out2)
end if
```

```
!$OMP END PARALLEL
```

- Important to distribute work evenly across threads (load balancing)

Imperial College
London

Simple parallel calculation

- Can use *sections* to assign tasks to threads:

```
!$OMP PARALLEL PRIVATE(threadID)
  NumThreads = omp_get_num_threads()
  threadID = omp_get_thread_num()
  !$OMP sections
    !$OMP section
      call subroutine1(in1,out1)
    !$OMP section
      call subroutine1(in2,out2)
  !$OMP end sections
!$OMP END PARALLEL
```

- Important to distribute work evenly across threads (load balancing)

Imperial College
London

Other useful OpenMP directives

- Single

```
!$OMP parallel
!$OMP single
  !code run by only one thread
!$OMP end single
!$OMP end parallel
```

- Used to run commands only once within parallel region
- Useful for: print statements, data input/output

Imperial College
London

Other useful OpenMP directives

- Single

```
!$OMP parallel
!$OMP single
  !code run by only one thread
!$OMP end single nowait ←
!$OMP end parallel
```

- Used to run commands only once within parallel region
- Useful for: print statements, data input/output
- Add *nowait* tag to allow other threads to continue while one thread is in *single* region

Imperial College
London

Synchronization

- Some threads may be given more work than others
- One thread may complete its tasks quickly and move very far ahead of the other threads
- **Barriers** keep the threads synchronized:

```
!$OMP parallel
!Some code
!$OMP barrier
!$OMP end parallel
```

- Threads will not continue past the barrier until all threads reach the barrier

Imperial College
London

Synchronization

- Some threads may be given more work than others
- One thread may complete its tasks quickly and move very far ahead of the other threads
- **Barriers** keep the threads synchronized:

```
!$OMP parallel
!Some code
!$OMP barrier
!$OMP end parallel
```

- Threads will not continue past the barrier until all threads reach the barrier
- There are *implicit* barriers at end of !\$OMP do and !\$OMP single blocks

Imperial College
London

Thread-safe routines

- What happens when you call sub-program from within parallel region?
- Each thread will call it's own "copy" of sub-program
 - All "local" variables declared within sub-program are private to thread

```
!$OMP parallel
call sub1(in1,in2,out1,out2)
!$OMP end parallel
!-----
subroutine sub1(in1,in2,out1,out2)
  use mod1
  implicit none
  real(kind=8) intent(in) :: in1,in2
  real(kind=8) intent(out) :: out1,out2
  real(kind=8) :: local1

  !should not modify mod1 variables
  !out1,out2 should (usually) be
  !private in the calling parallel region
end subroutine sub1
```

Imperial College
London


Basic questions:


1. Does code give same answer independent of the total number of threads?

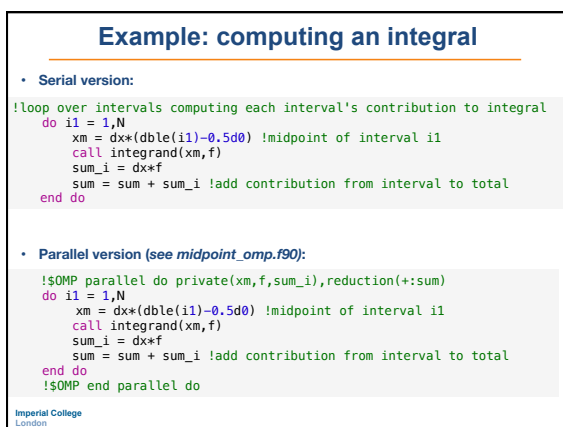
2. Is it independent of the order in which threads call the subroutine

If yes, the subroutine is *thread-safe*

Should not include OMP directives in subroutine called from within parallel region







Example: computing an integral

- Is there any actual performance gain?
 - Use `system_clock` and `omp_set_num_threads` (see `midpoint_time_omp.f90`)
- $N=1000$

| | | |
|-----------------------------|-------------------------|-----------------------------|
| <code>numThreads = 1</code> | <code>wall time=</code> | <code>2.30000005E-04</code> |
| <code>numThreads = 2</code> | <code>wall time=</code> | <code>6.97000010E-04</code> |
| <code>numThreads = 4</code> | <code>wall time=</code> | <code>1.09699997E-03</code> |
- Here, parallelization slows down the calculation! Why?
- Recall Amdahl's law, here $s > p$
- s/p will change as N increases...

Imperial College
London

Example: computing an integral

- Is there any actual performance gain?
 - Use `system_clock` and `omp_set_num_threads` (see `midpoint_time_omp.f90`)
- $N=1e7$

| | | |
|-----------------------------|-------------------------|---------------------------|
| <code>numThreads = 1</code> | <code>wall time=</code> | <code>1.11312997</code> |
| <code>numThreads = 2</code> | <code>wall time=</code> | <code>0.6055430174</code> |
| <code>numThreads = 4</code> | <code>wall time=</code> | <code>0.565499008</code> |
- Now, we see improved performance
- Speedup from two threads = 1.8
- No meaningful gain from four threads – laptop only has two cores

Imperial College
London
