

## Introduction to High Performance Scientific Computing

Autumn, 2016

Lecture 14

Imperial College  
London

Prasun Ray  
24 November 2016

## Vectorizing code

In general (Python, Fortran, Matlab,...), avoid for loops and *vectorize* calculations involving arrays.

Example:

```
In [27]: x=np.linspace(0,1,101)
```

```
In [28]: f = np.empty_like(x)
```

```
In [29]: for i in range(size(x)):
.....:     f[i] = cos(x[i])**2
.....:
```

```
In [30]: f = cos(x)**2    Vectorized version of loop
```

Imperial College  
London

## Vectorizing code

In general (Python, Fortran, Matlab,...), avoid for loops and *vectorize* calculations involving arrays.

Example:

```
In [27]: x=np.linspace(0,1,101)
```

```
In [28]: f = np.empty_like(x)
```

```
In [29]: for i in range(size(x)):
.....:     f[i] = cos(x[i])**2
.....:
```

```
In [30]: f = cos(x)**2    Vectorized version of loop
```

- Vectorized code will usually be faster, sometimes *much faster*

- Exception: parallelizing Fortran code with OpenMp:  
vectorized code → loops → parallel loops

Imperial College  
London

## Today

Programming example: from PDE  $\rightarrow$  algorithm  $\rightarrow$  serial code  $\rightarrow$  parallel code

Imperial College  
London

---

---

---

---

---

---

---

## Programming example

Task: Compute temperature distribution in a room

Imperial College  
London

---

---

---

---

---

---

---

## Programming example

Task: Compute temperature distribution in a room

Governing equation: Heat equation (diffusion equation):

$$\frac{\partial T}{\partial t} = \alpha \nabla^2 T + S(\mathbf{x}, t)$$

$$T(\mathbf{x}, t = 0) = f(\mathbf{x}) \quad \text{Initial condition}$$

Here,  $S$  is a *heat source*. Boundary conditions should also be specified as appropriate.

Problem: given the source, initial condition, and boundary conditions, solve for the temperature distribution,  $T(\mathbf{x}, t)$

Imperial College  
London

---

---

---

---

---

---

---

### Programming example

Today: 1-D problem

$$\frac{\partial T}{\partial t} = \frac{\partial^2 T}{\partial x^2} + S(x, t)$$

$$T(x, t = 0) = f(x)$$

Initial condition

$$T(x = 0, t) = a(t), \quad T(x = 1, t) = b(t)$$

Boundary conditions

$$0 \leq x \leq 1$$

Imperial College  
London

---

---

---

---

---

---

---

---

### Programming example

Today: 1-D problem

$$\frac{\partial T}{\partial t} = \frac{\partial^2 T}{\partial x^2} + S(x, t)$$

$$T(x, t = 0) = f(x)$$

Initial condition

$$T(x = 0, t) = a(t), \quad T(x = 1, t) = b(t)$$

Boundary conditions

$$0 \leq x \leq 1$$

First consider steady problem, e.g.,  $S = S(x)$ ,  $a$  and  $b$  are constants:

$$\frac{\partial^2 T}{\partial x^2} + S(x, t) = 0 \quad \text{Poisson equation}$$

Imperial College  
London

---

---

---

---

---

---

---

---

### Programming example

First consider steady problem, e.g.,  $S = S(x)$ ,  $a$  and  $b$  are constants:

$$\frac{\partial^2 T}{\partial x^2} + S(x) = 0$$

Notes:

1. This is an extremely simple problem, easy to write down the analytical solution
2. No need to use compiled language
3. Certainly no need to parallelize
4. But what about two-dimensional or three-dimensional problems?
  - Then, the picture changes considerably!
5. We are just considering the 1-D problem for illustrative purposes

Imperial College  
London

---

---

---

---

---

---

---

---

### Programming example

First consider steady problem, e.g.,  $S = S(x)$ ,  $a$  and  $b$  are constants:

$$\frac{\partial^2 T}{\partial x^2} + S(x) = 0 \quad \text{Poisson equation}$$

Numerical method:

1. Discretize the derivative:

$$\frac{\partial^2 T}{\partial x^2} \approx \frac{T_{i+1} - 2T_i + T_{i-1}}{\Delta x^2} \quad \text{2nd-order, centered scheme}$$

$$x_i = i * \Delta x, \quad i = 1, 2, \dots, N$$

$$(N + 1) * \Delta x = 1$$

Imperial College  
London

---

---

---

---

---

---

---

---

### Programming example

First consider steady problem, e.g.,  $S = S(x)$ ,  $a$  and  $b$  are constants:

$$\frac{\partial^2 T}{\partial x^2} + S(x) = 0 \quad \text{Poisson equation}$$

Numerical method:

1. Discretize the derivative:

$$\frac{\partial^2 T}{\partial x^2} \approx \frac{T_{i+1} - 2T_i + T_{i-1}}{\Delta x^2} \quad \text{2nd-order, centered scheme}$$

$$x_i = i * \Delta x, \quad i = 0, 1, 2, \dots, N + 1$$

$$(N + 1) * \Delta x = 1$$

With boundary conditions:  $T_0 = T_a, T_N = T_b$

Imperial College  
London

---

---

---

---

---

---

---

---

### Programming example

Equation for  $T_i$ :  $\frac{T_{i+1} - 2T_i + T_{i-1}}{\Delta x^2} = -S_i$

In matrix form:  $AT = b$

$$A = \begin{bmatrix} -2 & 1 & 0 & 0 & 0 & \dots & 0 \\ 1 & -2 & 1 & 0 & 0 & \dots & 0 \\ 0 & 1 & -2 & 1 & 0 & \dots & 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \vdots \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 1 & -2 & 1 \\ 0 & \dots & 0 & 0 & 0 & 1 & -2 \end{bmatrix}, \quad b = \frac{1}{\Delta x^2} \begin{bmatrix} -\Delta x^2 T_a - S_1 \\ -S_2 \\ \vdots \\ -S_i \\ \vdots \\ -S_{N-1} \\ -\Delta x^2 T_b - S_N \end{bmatrix}$$

- In 1-D, this is just a tridiagonal system of equations
- Easy to solve directly (with, say, *DGTSV*)

Imperial College  
London

---

---

---

---

---

---

---

---

### Programming example

- In two or three dimensions, A loses its simple banded structure
- Then, direct solution becomes very expensive for large N
- *Iterative* methods are a popular alternative

Imperial College  
London

---

---

---

---

---

---

---

---

### Programming example

- In two or three dimensions, A loses its simple banded structure
- Then, direct solution becomes very expensive for large N
- *Iterative* methods are a popular alternative
- Basic idea: rewrite  $Ax=b$  as  $A_1x = A_2x + b$
- Choose  $A_1$  so that it is easy to invert, then solve iterative system:
  - $A_1x^{k+1} = A_2x^k + b$ 
    - Requires guess,  $x^0$

Imperial College  
London

---

---

---

---

---

---

---

---

### Jacobi iteration

- Basic idea: rewrite  $Ax=b$  as  $A_1x = A_2x + b$
- Choose  $A_1$  so that it is easy to invert, then solve iterative system:
  - $A_1x^{k+1} = A_2x^k + b$ 
    - Requires guess,  $x^0$
  - *Jacobi iteration*: Choose  $A_1$  to be diagonal matrix (main diagonal of A):

$$\frac{T_{i+1}^{k-1} - 2T_i^k + T_{i-1}^{k-1}}{\Delta x^2} = -S_i$$

$$T_i^k = \frac{\Delta x^2}{2} S_i + \frac{1}{2} (T_{i+1}^{k-1} + T_{i-1}^{k-1})$$

Imperial College  
London

---

---

---

---

---

---

---

---

### Jacobi iteration

- Basic idea: rewrite  $Ax=b$  as  $A_1x = A_2x + b$
- Choose  $A_1$  so that it is easy to invert, then solve iterative system:
- $A_1x^{k+1} = A_2x^k + b$ 
  - Requires guess,  $x^0$
- **Jacobi iteration:** Choose  $A_1$  to be diagonal matrix (main diagonal of  $A$ ):

$$\frac{T_{i+1}^{k-1} - 2T_i^k + T_{i-1}^{k-1}}{\Delta x^2} = -S_i$$

$$T_i^k = \frac{\Delta x^2}{2} S_i + \frac{1}{2} (T_{i+1}^{k-1} + T_{i-1}^{k-1})$$

Main algorithm, easy to code!

Imperial College  
London

---

---

---

---

---

---

---

---

### Jacobi iteration in Fortran

- **Plan:**
  - Set parameters: a, b, n, tol
  - Construct grid  $x_i$
  - Construct source function,  $S(x)$ , initialize  $T=T(x,t=0)$
  - Iterate using formula below
    - Each iteration check if  $|T_k - T_{k-1}| < \text{tol}$

$$T_i^k = \frac{\Delta x^2}{2} S_i + \frac{1}{2} (T_{i+1}^{k-1} + T_{i-1}^{k-1})$$

Imperial College  
London

---

---

---

---

---

---

---

---

### Jacobi iteration in Fortran

- **One Fortran trick: set variables to be dimension(0:N+1)**
  - $x(0)=0, x(N+1)=1, T(0)=a, T(N+1)=b$
- **Then, easy to compute  $T_1$  using:**

$$T_i^k = \frac{\Delta x^2}{2} S_i + \frac{1}{2} (T_{i+1}^{k-1} + T_{i-1}^{k-1})$$

Imperial College  
London

---

---

---

---

---

---

---

---

### Jacobi iteration in Fortran

- One Fortran trick: set variables to be dimension(0:N+1)

- $x(0)=0$ ,  $x(N+1)=1$ ,  $T(0)=a$ ,  $T(N+1)=b$

- Then, easy to compute  $T_i$  using:

$$T_i^k = \frac{\Delta x^2}{2} S_i + \frac{1}{2} (T_{i+1}^{k-1} + T_{i-1}^{k-1})$$

Core part of code (see *jacobi1s.f90*):

```
do k1=1, kmax
  Tnew(1:n) = S(1:n)*dx2f + 0.5d0*(T(0:n-1) + T(2:n+1)) !Jacobi
  deltaT(k1) = maxval(abs(Tnew(1:n)-T(1:n))) !compute relative error
  T(1:n)=Tnew(1:n)      !update variable
  if (deltaT(k1)<tol) exit !check convergence criterion
end do
```

Imperial College  
London

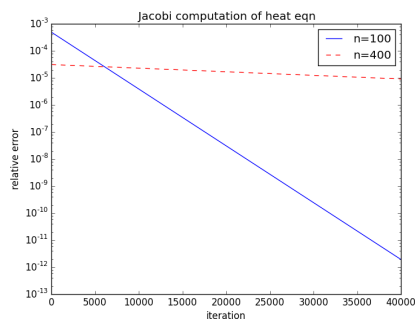
### Jacobi results

1. Does the solution converge?

Imperial College  
London


### Jacobi results

1. Does the solution converge? Yes, but very slowly for large n



### Jacobi results

1. Does it converge to the correct solution? Yes, can check that error  $\sim \Delta x^2$



### Jacobi results

- Jacobi is simplest, but *most inefficient* iterative solver
- Good illustration of basic ideas
- Better methods: Gauss-Seidel, SOR, conjugate gradient, multigrid

Imperial College  
London

### Parallel Jacobi

Let's now parallelize the solver with OpenMP

- Look for loops that can be parallelized
- Look for vectorized operations that can be converted to loops that can be parallelized

Serial:

```
Tnew(1:n) = S(1:n)*dx2f + 0.5d0*(T(0:n-1) + T(2:n+1)) !Jacobi
deltaT(k1) = maxval(abs(Tnew(1:n)-T(1:n))) !compute relative error
```

Imperial College  
London



## Parallel Jacobi

Let's now parallelize the solver with OpenMP

- Look for loops that can be parallelized
- Look for vectorized operations that can be converted to loops that can be parallelized

Parallel:

```
dmax=0.d0
!$omp parallel do reduction(max:dmax)
do i1=1,n
  Tnew(i1) = S(i1)*dx2f + 0.5d0*(T(i1-1) + T(i1+1))
  dmax = max(dmax,abs(Tnew(i1)-T(i1)))
end do
!$omp end parallel do
deltaT(k1) = dmax
```

Imperial College  
London

---

---

---

---

---

---

---

---

## Parallel Jacobi

Let's now parallelize the solver with OpenMP

- Look for loops that can be parallelized
- Look for vectorized operations that can be converted to loops that can be parallelized

Serial:

```
do i1=0,n+1
  x(i1) = i1*dx
end do
!-----
!set initial condition
T = (b-a)*x + a
!set source function
S = S0*sin(pi*x)
```

Imperial College  
London

---

---

---

---

---

---

---

---

## Parallel Jacobi

Let's now parallelize the solver with OpenMP

- Look for loops that can be parallelized
- Look for vectorized operations that can be converted to loops that can be parallelized

Parallel:

```
!$omp parallel do
do i1=0,n+1
  x(i1) = i1*dx
  T(i1) = (b-a)*x(i1) + a !set initial condition
  S(i1) = S0*sin(pi*x(i1)) !set source function
end do
!$omp end parallel do
```

Imperial College  
London

---

---

---

---

---

---

---

---

### Parallel Jacobi notes

- Will only see speedup with  $n > \sim 20000$  (commonly seen in 2D problems)
- See `jacobi1s_omp.f90`, `jacobi1_omp.py`
- **f2py and OpenMP:** `f2py -f90flags='-fopenmp' -lgomp -c jacobi1s_omp.f90 -m j1`

Imperial College  
London

---

---

---

---

---

---

---

---

### Time-dependent problem

Today: 1-D problem

$$\frac{\partial T}{\partial t} = \frac{\partial^2 T}{\partial x^2} + S(x, t)$$

$$T(x, t = 0) = f(x)$$

$$T(x = 0, t) = a(t), \quad T(x = 1, t) = b(t)$$

$$0 \leq x \leq 1$$

Simple (inefficient) approach: *method of lines*

1. Discretize spatial variable  $\rightarrow N+2$  points between 0 and 1
2. Solve resulting N ODEs with solver of choice (*odeint*, *ode15s*, ...)

Imperial College  
London

---

---

---

---

---

---

---

---

### Time-dependent problem

Again, we discretize the derivative:

$$\frac{\partial^2 T}{\partial x^2} \approx \frac{T_{i+1} - 2T_i + T_{i-1}}{\Delta x^2}$$

$$x_i = i * \Delta x, \quad i = 0, 1, 2, \dots, N+1$$

$$(N+1) * \Delta x = 1$$

Imperial College  
London

---

---

---

---

---

---

---

---

### Time-dependent problem

Again, we discretize the derivative as:

$$\frac{\partial^2 T}{\partial x^2} \approx \frac{T_{i+1} - 2T_i + T_{i-1}}{\Delta x^2}$$

$$x_i = i * \Delta x, \quad i = 0, 1, 2, \dots, N+1$$

$$(N+1) * \Delta x = 1$$

So, we have  $N$  ODEs:

$$\frac{dT_i}{dt} = S_i(t) + \frac{T_{i+1} - 2T_i + T_{i-1}}{\Delta x^2}, \quad i = 1, 2, \dots, N$$

with the boundary conditions substituted in the RHS when needed.

Imperial College  
London

---

---

---

---

---

---

---

---

### Solving single ODE in python (lecture 6)

- Use `odeint` from `scipy.integrate` module to solve:

$$\frac{dy}{dt} = -ay$$

- **Basic idea:** discretize time,  $t = 0, dt, \dots, N*dt$ , and starting from  $y(0)$  march forward in time and compute  $y(dt), \dots, y(N*dt)$
- `odeint` chooses the stepsize,  $dt$ , so that error tolerances are satisfied
- Need to specify:
  - Initial condition
  - Timespan for integration
  - A Python function which provides RHS of the ODE to `odeint`
- Look at `ode_example.py` and lab 4

Imperial College  
London

---

---

---

---

---

---

---

---

### Time-dependent problem

Solving  $N$  ODEs:

$$\frac{dT_i}{dt} = S_i(t) + \frac{T_{i+1} - 2T_i + T_{i-1}}{\Delta x^2}, \quad i = 1, 2, \dots, N$$

- Will need to provide  $N$  initial conditions when calling `odeint`.
- The python function which provides RHS to `odeint` will:
  - Take  $t$  and  $T_0, \dots, T_N$  and any other needed parameters as input
  - Return  $N$  values for  $dT/dt$  as output
- No need for Fortran for 1D problems, but may be faster for two and three dimensions.

Imperial College  
London

---

---

---

---

---

---

---

---