


Introduction to High Performance Scientific Computing

Autumn, 2016

Lecture 15


Parallel computing paradigms



Distributed memory

- Each (4-core) chip has its own memory
- The chips are connected by network 'cables'
- MPI coordinates communication between two or more CPUs

Parallel computing paradigms



Related approaches:

- Hybrid programming: mix of shared-memory (OpenMP) and distributed-memory (MPI) programming
- GPU's: Shared memory programming (CUDA or OpenCL)
- Coprocessors and co-array programming

MPI intro

- **MPI: *Message Passing Interface***
- **Standard for exchanging data between processors**
- **Supports Fortran, c, C++**
- **Can also be used with Python**


OpenMP schematic

Program starts with
single *master thread*

Then, launch parallel
region with multiple
threads.

Each thread has
access to all
variables introduced
previously

Can end parallel
region if/when
desired and launch
parallel regions again
in future as needed




MPI schematic

Program starts with
all processes
running

MPI controls
communication
between processes

Start program



Parallel region (4 processes)


MPI intro

- Basic idea: calls to MPI subroutines control data exchange between processors
- Example:

```
call MPI_BCAST(n, 1, MPI_INTEGER, 0, MPI_COMM_WORLD, ierr)
```

This will send the integer `n` which has size `1` from processor `0` to all of the other processors.

MPI broadcast



MPI intro

- Basic idea: calls to MPI subroutines control data exchange between processors
- Example:

```
call MPI_BCAST(n, 1, MPI_INTEGER, 0, MPI_COMM_WORLD, ierr)
```

This will send the integer `n` which has size `1` from processor `0` to all of the other processors.

Generally, need to specify:

- source and/or destination of message
- size of data contained in message
- type of data contained in message (integer, double precision, ...)
- the data itself (or its location)

Fortran code structure

! Basic Fortran 90 code structure

!1. Header

```
program template
```

!2. Variable declarations (e.g. integers, real numbers,...)

!3. basic code: input, loops, if-statements, subroutine calls

```
print *, 'template code'
```

!4. End program

```
end program template
```

! To compile this code:

```
! $ gfortran -o f90template.exe f90template.f90
```

```
! To run the resulting executable: $ ./f90template.exe
```

MPI intro

! Basic MPI + Fortran 90 code structure

See *mpif90template.f90*

!1. Header

```
program template
```

```
  use mpi
```

!2a. Variable declarations (e.g. integers, real numbers,...)

```
  integer :: myid, numprocs, ierr
```

!2b. Initialize MPI

```
  call MPI_INIT(ierr)
```

```
  call MPI_COMM_RANK(MPI_COMM_WORLD, myid, ierr)
```

```
  call MPI_COMM_SIZE(MPI_COMM_WORLD, numprocs, ierr)
```

!3. basic code: input, loops, if-statements, subroutine calls

```
  print *, 'this is proc # ',myid, 'of ', numprocs
```

!4. End program

```
  call MPI_FINALIZE(ierr)
```

```
end program template
```

! To compile this code:

```
! $ mpif90 -o mpitemplate.exe mpif90template.f90
```

! To run the resulting executable with 4 processes:\$ mpiexec -n 4 mpitemplate.exe

MPI intro

! Basic MPI + Fortran 90 code structure

See *mpif90template.f90*

!1. Header

```
program template
```

```
  use mpi
```

!2a. Variable declarations (e.g. integers, real numbers,...)

```
  integer :: myid, numprocs, ierr
```

!2b. Initialize MPI

```
  call MPI_INIT(ierr)
```

```
  call MPI_COMM_RANK(MPI_COMM_WORLD, myid, ierr)
```

```
  call MPI_COMM_SIZE(MPI_COMM_WORLD, numprocs, ierr)
```

!3. basic code: input, loops, if-statements, subroutine calls

```
  print *, 'this is proc # ',myid, 'of ', numprocs
```

!4. End program

```
  call MPI_FINALIZE(ierr)
```

```
end program template
```

! To compile this code:

```
! $ mpif90 -o mpitemplate.exe mpif90template.f90
```

! To run the resulting executable with 4 processes:\$ mpiexec -n 4 mpitemplate.exe

MPI intro

! Basic MPI + Fortran 90 code structure

See *mpif90template.f90*

!1. Header

```
program template
```

```
  use mpi
```

!2a. Variable declarations (e.g. integers, real numbers,...)

```
  integer :: myid, numprocs, ierr
```

!2b. Initialize MPI

```
  call MPI_INIT(ierr)
```

```
  call MPI_COMM_RANK(MPI_COMM_WORLD, myid, ierr)
```

```
  call MPI_COMM_SIZE(MPI_COMM_WORLD, numprocs, ierr)
```

!3. basic code: input, loops, if-statements, subroutine calls

```
  print *, 'this is proc # ',myid, 'of ', numprocs
```

!4. End program

```
  call MPI_FINALIZE(ierr)
```

```
end program template
```

! To compile this code:

```
! $ mpif90 -o mpitemplate.exe mpif90template.f90
```

```
! To run the resulting executable with 4 processes:$ mpiexec -n 4 mpitemplate.exe
```

MPI intro

- **Compile + run:**


```
$ mpif90 -o mpif90template.exe mpif90template.f90
```

```
$ mpiexec -n 4 mpif90template.exe
```

```
this is proc #          0 of          4
this is proc #          3 of          4
this is proc #          1 of          4
this is proc #          2 of          4
```

Note: The number of processes specified with `mpiexec` can be larger than the number of cores on your machine, but then tasks are run sequentially.

MPI+Fortran example: computing an integral



- Estimate integral with midpoint rule,

$$I = \int_0^1 \frac{4}{1+x^2} dx$$

MPI+Fortran quadrature

Two most important tasks:

1. Decide how many intervals per processor
2. Each processor will compute its own partial sum, `sum_proc`, how do we compute `sum(sum_proc)`?

MPI+Fortran quadrature

Two most important tasks:

1. **Decide how many intervals per processor**
2. **Each processor will compute its own partial sum, `sum_proc`, how do we compute `sum(sum_proc)`?**
 - **`N` = number of intervals**
 - **`numprocs` = number of processors**
 - **Need to compute `Nper_proc`: intervals per processor**

MPI+Fortran quadrature

- **N = number of intervals**
- **numprocs = number of processors**
- **Need to compute Nper_proc: intervals per processor**
 - **Basic idea: if** $N = 8 * \text{numprocs}$, $N_{\text{per_proc}} = 8$
 - **But, if** $N \leq \text{numprocs}$, **$N/\text{numprocs} = 0$**

$$N_{\text{per_proc}} = (N + \text{numprocs} - 1) / \text{numprocs}$$


MPI+Fortran quadrature

Two most important tasks:

1. Decide how many intervals per processor
2. Each processor will compute its own partial sum, `sum_proc`, how do we compute `sum(sum_proc)`?

Use `MPI_REDUCE`

MPI reduce



MPI+Fortran quadrature

Two most important tasks:

1. Decide how many intervals per processor
2. Each processor will compute its own partial sum, `sum_proc`, how do we compute `sum(sum_proc)`?
 - Use `MPI_REDUCE`
 - Reduction options: `MPI_MAX`, `MPI_MIN`, `MPI_SUM`, `MPI_PROD`

MPI+Fortran quadrature

Two most important tasks:

1. Decide how many intervals per processor
2. Each processor will compute its own partial sum, `sum_proc`, how do we compute `sum(sum_proc)`?
 - Use `MPI_REDUCE`
 - Reduction options: `MPI_MAX`, `MPI_MIN`, `MPI_SUM`, `MPI_PROD`
 - For quadrature, we need `MPI_SUM`

MPI+Fortran quadrature

For quadrature, we need MPI_SUM:

```
call MPI_REDUCE(data, result, 1, MPI_DOUBLE_PRECISION,  
               MPI_SUM, 0, MPI_COMM_WORLD, ierr)
```

This will:

1. Collect the double precision variable `data` which has size `1` from each processor.
2. Compute the sum (because we have chosen `MPI_SUM`) and store the value in `result` on processor `0`.

Note: Only processor 0 will have the final sum. With `MPI_ALLREDUCE`, the result will be on every processor.

MPI+Fortran quadrature

midpoint_p.f90: distribute data

```
!set number of intervals per processor
```

```
  Nper_proc = (N + numprocs - 1)/numprocs
```

```
!starting and ending points for processor
```

```
  istart = myid * Nper_proc + 1
```

```
  iend = (myid+1) * Nper_proc
```

```
  if (iend>N) iend = N
```


MPI+Fortran quadrature

midpoint_p.f90: 1. distribute data, 2. compute sum_proc

```
!set number of intervals per processor
```

```
Nper_proc = (N + numprocs - 1)/numprocs
```

```
!starting and ending points for processor
```

```
istart = myid * Nper_proc + 1
```

```
iend = (myid+1) * Nper_proc
```

```
if (iend>N) iend = N
```

```
!loop over intervals computing each interval's contribution to  
integral
```

```
do i1 = istart,iend
```

```
  xm = dx*(i1-0.5) !midpoint of interval i1
```

```
  call integrand(xm,f)
```

```
  sum_i = dx*f
```

```
  sum_proc = sum_proc + sum_i !add contribution from interval  
to total integral
```

```
end do
```

MPI+Fortran quadrature

midpoint_p.f90: 1. distribute data, 2. compute sum_proc, 3. reduction

```
!set number of intervals per processor
```

```
Nper_proc = (N + numprocs - 1)/numprocs
```

```
!starting and ending points for processor
```

```
istart = myid * Nper_proc + 1
```

```
iend = (myid+1) * Nper_proc
```

```
if (iend>N) iend = N
```

```
!loop over intervals computing each interval's contribution to integral
```

```
do i1 = istart,iend
```

```
  xm = dx*(i1-0.5) !midpoint of interval i1
```

```
  call integrand(xm,f)
```

```
  sum_i = dx*f
```

```
  sum_proc = sum_proc + sum_i !add contribution from interval to
```

```
total integral
```

```
end do
```

```
!collect double precision variable, sum, with size 1 on process 0 using  
the MPI_SUM option
```

```
call MPI_REDUCE(sum_proc,sum,1,MPI_DOUBLE_PRECISION,MPI_SUM,
```

```
0,MPI_COMM_WORLD,ierr)
```

MPI+Fortran quadrature

Compile and run:

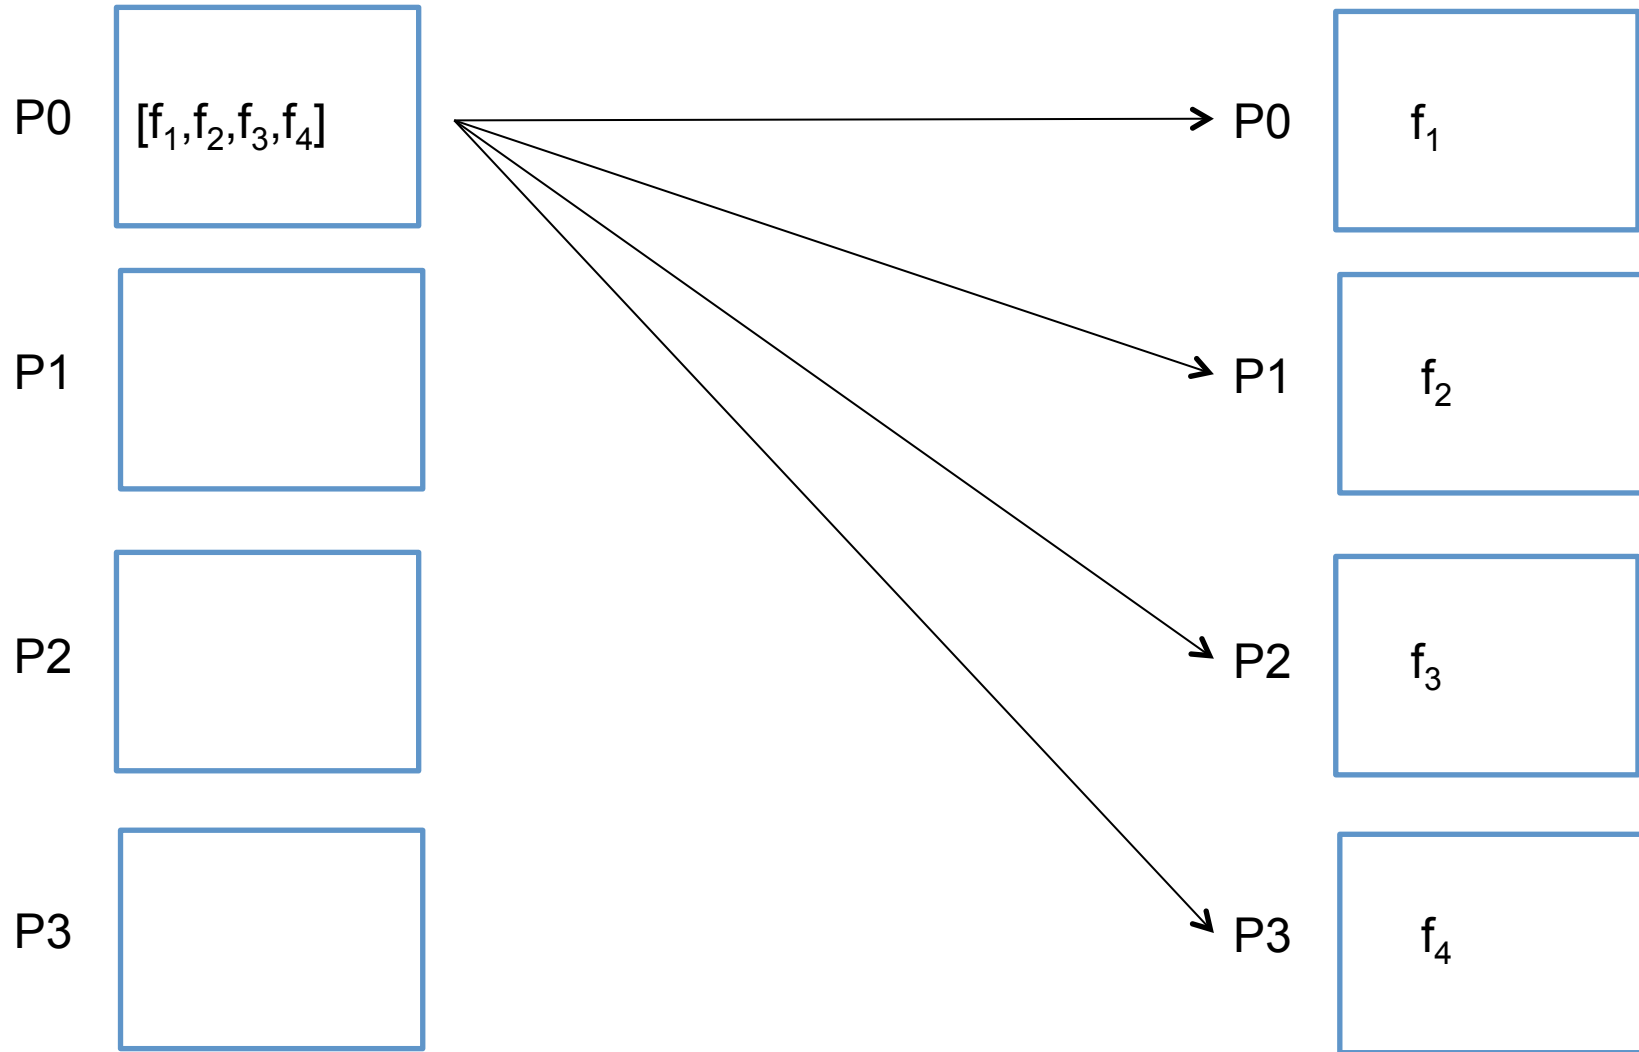
```
$ mpif90 -o midpoint_p.exe midpoint_p.f90

$ mpiexec -n 2 midpoint_p.exe
number of intervals =      1000
number of procs =      2
Nper_proc=      500
The partial sum on proc #      0 is:      1.8545905426699112
The partial sum on proc #      1 is:      1.2870021942532193
N=      1000
sum=      3.1415927369231307
error=      8.3333337563828991E-008
```


Other collective operations

- Scatter and gather

MPI scatter



MPI gather



Other collective operations

- **Scatter and gather**
 - ***Gather* all particles on processor**
 - **Compute interaction forces for particles on that processor**

$$\frac{d^2 \mathbf{x}_i}{dt^2} = \sum_{j=1}^N f(|\mathbf{x}_i - \mathbf{x}_j|), \quad i = 1, 2, \dots, N$$

- **Avoid for big problems (why?)**

MPI collective data movement

