

## Introduction to High Performance Scientific Computing

Autumn, 2016

Lecture 15

Imperial College  
London

Prasun Ray  
28 November 2016

---

---

---


---

---

---

---

### Parallel computing paradigms



#### Distributed memory

- Each (4-core) chip has its own memory
- The chips are connected by network 'cables'
- MPI coordinates communication between two or more CPUs

Imperial College  
London

---

---

---


---

---

---

---

### Parallel computing paradigms



#### Related approaches:

- Hybrid programming: mix of shared-memory (OpenMP) and distributed-memory (MPI) programming
- GPU's: Shared memory programming (CUDA or OpenCL)
- Coprocessors and co-array programming

Imperial College  
London

---

---

---

---

---

---

---

### MPI intro

- **MPI: Message Passing Interface**
- Standard for exchanging data between processors
- Supports Fortran, c, C++
- Can also be used with Python

Imperial College  
London

---

---

---

---

---

---

---


### OpenMP schematic

Program starts with  
single *master thread*

Then, launch parallel  
region with multiple  
threads.

Each thread has  
access to all  
variables introduced  
previously

Can end parallel  
region if/when  
desired and launch  
parallel regions again  
in future as needed



Imperial College  
London

---

---

---

---

---


---

---

### MPI schematic

Program starts with  
all processes  
running

MPI controls  
communication  
between processes



Imperial College  
London

---

---

---

---

---

---

---

### MPI intro

- Basic idea: calls to MPI subroutines control data exchange between processors
- Example:  
call MPI\_BCAST(*n*, 1, MPI\_INTEGER, 0, MPI\_COMM\_WORLD, ierr)

This will send the integer *n* which has size 1 from processor 0 to all of the other processors.

Imperial College  
London

---

---

---

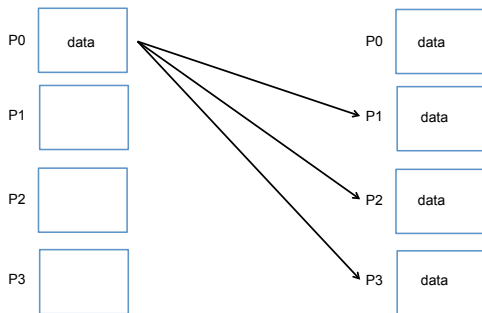
---

---

---

---

### MPI broadcast



Imperial College  
London

---

---

---

---

---

---

---

### MPI intro

- Basic idea: calls to MPI subroutines control data exchange between processors
- Example:  
call MPI\_BCAST(*n*, 1, MPI\_INTEGER, 0, MPI\_COMM\_WORLD, ierr)

This will send the integer *n* which has size 1 from processor 0 to all of the other processors.

Generally, need to specify:

- source and/or destination of message
- size of data contained in message
- type of data contained in message (integer, double precision, ...)
- the data itself (or its location)

Imperial College  
London

---

---

---

---

---

---

---

## Fortran code structure

! Basic Fortran 90 code structure

!1. Header

program template

!2. Variable declarations (e.g. integers, real numbers,...)

!3. basic code: input, loops, if-statements, subroutine calls  
print \*, 'template code'

!4. End program

end program template

! To compile this code:

! \$ gfortran -o f90template.exe f90template.f90

! To run the resulting executable: \$ ./f90template.exe

Imperial College  
London

## MPI intro

! Basic MPI + Fortran 90 code structure **See mpif90template.f90**

!1. Header

program template

use mpi

!2a. Variable declarations (e.g. integers, real numbers,...)  
integer :: myid, numprocs, ierr

!2b. Initialize MPI

call MPI\_INIT(ierr)

call MPI\_COMM\_RANK(MPI\_COMM\_WORLD, myid, ierr)

call MPI\_COMM\_SIZE(MPI\_COMM\_WORLD, numprocs, ierr)

!3. basic code: input, loops, if-statements, subroutine calls  
print \*, 'this is proc # ',myid, 'of ', numprocs

!4. End program

call MPI\_FINALIZE(ierr)

end program template

! To compile this code:

! \$ mpif90 -o mpitemplate.exe mpif90template.f90

! To run the resulting executable with 4 processes: \$ mpiexec -n 4 mpitemplate.exe

Imperial College  
London

## MPI intro

! Basic MPI + Fortran 90 code structure **See mpif90template.f90**

!1. Header

program template

use mpi

!2a. Variable declarations (e.g. integers, real numbers,...)  
integer :: myid, numprocs, ierr

!2b. Initialize MPI

call MPI\_INIT(ierr)

call MPI\_COMM\_RANK(MPI\_COMM\_WORLD, myid, ierr)

call MPI\_COMM\_SIZE(MPI\_COMM\_WORLD, numprocs, ierr)

!3. basic code: input, loops, if-statements, subroutine calls  
print \*, 'this is proc # ',myid, 'of ', numprocs

!4. End program

call MPI\_FINALIZE(ierr)

end program template

! To compile this code:

! \$ mpif90 -o mpitemplate.exe mpif90template.f90

! To run the resulting executable with 4 processes: \$ mpiexec -n 4 mpitemplate.exe

Imperial College  
London

## MPI intro

! Basic MPI + Fortran 90 code structure **See mpif90template.f90**

```

!1. Header
program template
  use mpi

  !2a. Variable declarations (e.g. integers, real numbers,...)
  integer :: myid, numprocs, ierr

  !2b. Initialize MPI
  call MPI_INIT(ierr)
  call MPI_COMM_RANK(MPI_COMM_WORLD, myid, ierr)
  call MPI_COMM_SIZE(MPI_COMM_WORLD, numprocs, ierr)

  !3. basic code: input, loops, if-statements, subroutine calls
  print *, 'this is proc # ',myid, ' of ', numprocs

!4. End program
  call MPI_FINALIZE(ierr)
end program template

! To compile this code:
! $ mpif90 -o mpitemplate.exe mpif90template.f90
! To run the resulting executable with 4 processes: $ mpiexec -n 4 mpitemplate.exe

```

Imperial College  
London

## MPI intro

- Compile + run:

```
$ mpif90 -o mpif90template.exe mpif90template.f90
```

```
$ mpiexec -n 4 mpif90template.exe
```

```


this is proc #      0 of      4
this is proc #      3 of      4
this is proc #      1 of      4
this is proc #      2 of      4

```

**Note:** The number of processes specified with mpiexec can be larger than the number of cores on your machine, but then tasks are run sequentially.

Imperial College  
London

## MPI+Fortran example: computing an integral



- Estimate integral with midpoint rule,

$$I = \int_0^1 \frac{4}{1+x^2} dx$$

Imperial College  
London

### MPI+Fortran quadrature

Two most important tasks:

1. Decide how many intervals per processor
2. Each processor will compute its own partial sum, `sum_proc`, how do we compute `sum(sum_proc)`?

Imperial College  
London

---

---

---

---

---

---

---

---

### MPI+Fortran quadrature

Two most important tasks:

1. Decide how many intervals per processor
2. Each processor will compute its own partial sum, `sum_proc`, how do we compute `sum(sum_proc)`?
  - `N` = number of intervals
  - `numprocs` = number of processors
  - Need to compute `Nper_proc`: intervals per processor

Imperial College  
London

---

---

---

---

---

---

---

---

### MPI+Fortran quadrature

- `N` = number of intervals
- `numprocs` = number of processors
- Need to compute `Nper_proc`: intervals per processor
- Basic idea: if `N = 8 * numprocs`, `Nper_proc = 8`
- But, if `N <= numprocs`, `N/numprocs = 0`

$$Nper\_proc = (N + numprocs - 1) / numprocs$$

Imperial College  
London

---

---

---

---

---

---

---

---

### MPI+Fortran quadrature

---

Two most important tasks:

1. Decide how many intervals per processor
2. Each processor will compute its own partial sum, `sum_proc`, how do we compute `sum(sum_proc)`?

Use `MPI_REDUCE`

Imperial College  
London

---

---

---

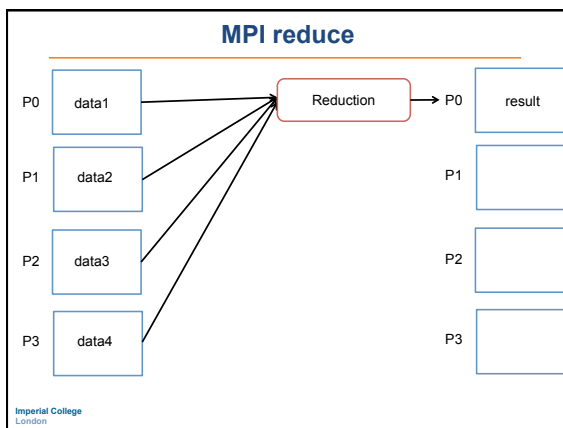
---

---

---

---

---




---

---

---

---

---

---

---

---

### MPI+Fortran quadrature

---

Two most important tasks:

1. Decide how many intervals per processor
2. Each processor will compute its own partial sum, `sum_proc`, how do we compute `sum(sum_proc)`?

- Use `MPI_REDUCE`
- Reduction options: `MPI_MAX`, `MPI_MIN`, `MPI_SUM`, `MPI_PROD`

Imperial College  
London

---

---

---

---

---

---

---

---

## MPI+Fortran quadrature

Two most important tasks:

1. Decide how many intervals per processor
2. Each processor will compute its own partial sum, `sum_proc`, how do we compute `sum(sum_proc)`?
  - Use `MPI_REDUCE`
  - Reduction options: `MPI_MAX`, `MPI_MIN`, `MPI_SUM`, `MPI_PROD`
  - For quadrature, we need `MPI_SUM`

Imperial College  
London

---

---

---

---

---

---

---

---

## MPI+Fortran quadrature

For quadrature, we need `MPI_SUM`:

call `MPI_REDUCE(data, result, 1, MPI_DOUBLE_PRECISION, MPI_SUM, 0, MPI_COMM_WORLD, ierr)`

This will:

1. Collect the double precision variable `data` which has size 1 from each processor.
2. Compute the sum (because we have chosen `MPI_SUM`) and store the value in `result` on processor 0.

Note: Only processor 0 will have the final sum. With `MPI_ALLREDUCE`, the result will be on every processor.

Imperial College  
London

---

---

---

---

---

---

---

---

## MPI+Fortran quadrature

*midpoint\_p.f90*: distribute data

```
!set number of intervals per processor
Nper_proc = (N + numprocs - 1)/numprocs

!starting and ending points for processor
istart = myid * Nper_proc + 1
iend = (myid+1) * Nper_proc
if (iend>N) iend = N
```

Imperial College  
London

---

---

---

---

---

---

---

---



### MPI+Fortran quadrature

*midpoint\_p.f90*: 1. distribute data, 2. compute sum\_proc

```
!set number of intervals per processor
Nper_proc = (N + numprocs - 1)/numprocs

!starting and ending points for processor
istart = myid * Nper_proc + 1
iend = (myid+1) * Nper_proc
if (iend-N) iend = N

!loop over intervals computing each interval's contribution to
integral
do i1 = istart,iend
  xm = dx*(i1-0.5) !midpoint of interval i1
  call integrand(xm,f)
  sum_i = dx*f
  sum_proc = sum_proc + sum_i !add contribution from interval
to total integral
end do
```

Imperial College  
London

---

---

---

---

---

---

---

---

### MPI+Fortran quadrature

*midpoint\_p.f90*: 1. distribute data, 2. compute sum\_proc, 3. reduction

```
!set number of intervals per processor
Nper_proc = (N + numprocs - 1)/numprocs

!starting and ending points for processor
istart = myid * Nper_proc + 1
iend = (myid+1) * Nper_proc
if (iend-N) iend = N

!loop over intervals computing each interval's contribution to integral
do i1 = istart,iend
  xm = dx*(i1-0.5) !midpoint of interval i1
  call integrand(xm,f)
  sum_i = dx*f
  sum_proc = sum_proc + sum_i !add contribution from interval to
total integral
end do
!collect double precision variable, sum, with size 1 on process 0 using
the MPI_SUM option
call MPI_REDUCE(sum_proc,sum,1,MPI_DOUBLE_PRECISION,MPI_SUM,
0,MPI_COMM_WORLD,ierr)
```

Imperial College  
London

---

---

---

---

---

---

---

---

### MPI+Fortran quadrature

Compile and run:

```
$ mpif90 -o midpoint_p.exe midpoint_p.f90

$ mpiexec -n 2 midpoint_p.exe
number of intervals = 1000
number of procs = 2
Nper_proc= 500
The partial sum on proc # 0 is: 1.8545905426699112
The partial sum on proc # 1 is: 1.2870021942532193
N= 1000
sum= 3.1415927369231307
error= 8.3333337563828991E-008
```

Imperial College  
London

---

---

---

---

---

---

---

---

### Other collective operations

- Scatter and gather

Imperial College  
London

---

---

---


---

---

---

---

### MPI scatter



Imperial College  
London

---

---

---


---

---

---

---

### MPI gather



Imperial College  
London

---

---

---

---

---

---

---

### Other collective operations

- Scatter and gather
  - Gather** all particles on processor
  - Compute interaction forces for particles on that processor

$$\frac{d^2 \mathbf{x}_i}{dt^2} = \sum_{j=1}^N f(|\mathbf{x}_i - \mathbf{x}_j|), \quad i = 1, 2, \dots, N$$

- Avoid for big problems (why?)

Imperial College  
London

---

---

---

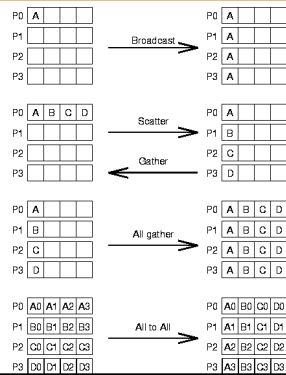
---

---

---

---

### MPI collective data movement



Imperial College  
London

From Using MPI

---

---

---

---

---

---

---