

Introduction to High Performance Scientific Computing

Autumn, 2016

Lecture 17


Project Part 1

- Random walks in 1-D
- Consider equations of the form:
$$X(t+dt) = X(t) + F[X(t),dt]$$
- Here, X and F , are both random variables
- Simplest example: $F = \pm dx$ based on flip of a coin \rightarrow random walk
- Compute several *realizations*, $X_1(t)$, $X_2(t)$,... $X_M(t)$ and then compute statistics over these M realizations
- These statistics (the mean and stdev) will generally depend on time
- Random walk on a network follows same idea

Send/Recv and domain decomposition

A parallel computation computes a potential field, $f(x,y,z,t)$ on four processors.

P0, P1, P2, P3 solve for f in separate subdomains



- How would you compute the gradient?


$$\nabla f = \left(\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z} \right)$$

- No problems with x and z directions
- But what about y?

Last time: Send/Recv

A parallel computation computes a potential field, $f(x,y,z,t)$ on four processors.

P0, P1, P2, P3 solve for f in separate subdomains



- Send and Recv *must* be paired:
 - If yellow sends to blue, blue must recv from yellow
- Yellow and blue should not send to each other simultaneously
- Instead:
Yellow → Blue
Blue → Green
Green → Red
Red → Yellow (if periodic)

and then the reverse

Parallel differentiation example

```
!-----  
!Send data at top boundary up to next processor  
!i.e. send f(nlocal+1) to myid+1 and store it there as f(1)  
!data from myid=numprocs-1 is sent to myid=0  
!-----
```

```
if (myid<numprocs-1) then  
    receiver = myid+1
```

```
else  
    receiver = 0
```

```
end if
```

```
if (myid>0) then  
    sender = myid-1
```

```
else  
    sender = numprocs-1
```

```
end if
```

```
call MPI_SEND(f(Nlocal+1),1,MPI_DOUBLE_PRECISION,receiver,0,  
              MPI_COMM_WORLD,ierr)
```

```
call MPI_RECV(f(1) 1,MPI_DOUBLE_PRECISION,sender,MPI_ANY_TAG,  
              MPI_COMM_WORLD,status,ierr)
```

Today

Solving the (steady) 2D heat equation

Heat equation

Task: Compute temperature distribution in a room

Governing equation: Heat equation (diffusion equation):

$$\frac{\partial T}{\partial t} = \alpha \nabla^2 T + S(\mathbf{x}, t)$$

$$T(\mathbf{x}, t = 0) = f(\mathbf{x}) \quad \text{Initial condition}$$

Here, S is a *heat source*. Boundary conditions should also be specified as appropriate.

Problem: given the source, initial condition, and boundary conditions, solve for the temperature distribution, $T(\mathbf{x}, t)$

1-D (steady) heat equation

First consider steady problem, e.g., $S = S(x)$, a and b are constants:

$$\frac{\partial^2 T}{\partial x^2} + S(x) = 0 \quad \text{Poisson equation}$$

Numerical method:

1. Discretize the derivative:

$$\frac{\partial^2 T}{\partial x^2} \approx \frac{T_{i+1} - 2T_i + T_{i-1}}{\Delta x^2} \quad \text{2nd-order, centered scheme}$$

$$x_i = i * \Delta x, \quad i = 0, 1, 2, \dots, N + 1$$

$$(N + 1) * \Delta x = 1$$

With boundary conditions: $T_0 = T_a, \quad T_N = T_b$

Programming example

Equation for T_i :
$$\frac{T_{i+1} - 2T_i + T_{i-1}}{\Delta x^2} = -S_i$$

In matrix form: $AT = b$

$$A = \begin{bmatrix} -2 & 1 & 0 & 0 & 0 & \dots & 0 \\ 1 & -2 & 1 & 0 & 0 & \dots & 0 \\ 0 & 1 & -2 & 1 & 0 & \dots & 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots & \vdots & \vdots \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 1 & -2 & 1 \\ 0 & \dots & 0 & 0 & 0 & 1 & -2 \end{bmatrix}, \quad b = \frac{1}{\Delta x^2} \begin{bmatrix} -\Delta x^2 T_a - S_1 \\ -S_2 \\ \vdots \\ -S_i \\ \vdots \\ -S_{N-1} \\ -\Delta x^2 T_b - S_N \end{bmatrix}$$

- In 1-D, this is just a tridiagonal system of equations
- Easy to solve directly (with, say, *DGTSV*)

Jacobi iteration

- **Basic idea:** rewrite $Ax=b$ as $A_1x = A_2x + b$
- **Choose A_1** so that it is easy to invert, then solve iterative system:
- $A_1x^{k+1} = A_2x^k + b$
 - Requires guess, x^0
- **Jacobi iteration:** Choose A_1 to be diagonal matrix (main diagonal of A):

$$\frac{T_{i+1}^{k-1} - 2T_i^k + T_{i-1}^{k-1}}{\Delta x^2} = -S_i$$

$$T_i^k = \frac{\Delta x^2}{2} S_i + \frac{1}{2} (T_{i+1}^{k-1} + T_{i-1}^{k-1})$$

Main algorithm, easy to code!

Jacobi iteration in Fortran

- **One Fortran trick: set variables to be dimension(0:N+1)**
 - $x(0)=0$, $x(N+1)=1$, $T(0)=a$, $T(N+1)=b$
 - **Then, easy to compute T_1 using:**

$$T_i^k = \frac{\Delta x^2}{2} S_i + \frac{1}{2} (T_{i+1}^{k-1} + T_{i-1}^{k-1})$$

Core part of code (see *jacobi1s.f90*):

```
do k1=1, kmax
```

```
  Tnew(1:n) = S(1:n)*dx2f + 0.5d0*(T(0:n-1) + T(2:n+1)) !Jacobi
```

```
  deltaT(k1) = maxval(abs(Tnew(1:n)-T(1:n))) !compute relative error
```

```
  T(1:n)=Tnew(1:n) !update variable
```

```
  if (deltaT(k1)<tol) exit !check convergence criterion
```

```
end do
```

Parallel Jacobi

Let's now parallelize the solver with OpenMP

- Look for loops that can be parallelized
- Look for vectorized operations that can be converted to loops that can be parallelized

Parallel:

```
    dmax=0.d0
    !$omp parallel do reduction(max:dmax)
    do i1=1,n
        Tnew(i1) = S(i1)*dx2f + 0.5d0*(T(i1-1) + T(i1+1))
        dmax = max(dmax,abs(Tnew(i1)-T(i1)))
    end do
    !$omp end parallel do
    deltaT(k1) = dmax
```

2-D (steady) heat equation

$$\frac{\partial^2 T}{\partial x^2} + \frac{\partial^2 T}{\partial y^2} + S(x, y) = 0 \quad \text{Poisson equation}$$

2-D (steady) heat equation

$$\frac{\partial^2 T}{\partial x^2} + \frac{\partial^2 T}{\partial y^2} + S(x, y) = 0 \quad \text{Poisson equation}$$

Numerical method:

1. Discretize the derivative:

$$\left(\frac{\partial^2 T}{\partial x^2} \right)_{i,j} \approx \frac{T_{i+1,j} - 2T_{i,j} + T_{i-1,j}}{\Delta x^2}$$

$$\left(\frac{\partial^2 T}{\partial y^2} \right)_{i,j} \approx \frac{T_{i,j+1} - 2T_{i,j} + T_{i,j-1}}{\Delta y^2}$$

$$\Delta x = \Delta y = \Delta$$

2nd-order, centered scheme

2-D (steady) heat equation

$$\frac{\partial^2 T}{\partial x^2} + \frac{\partial^2 T}{\partial y^2} + S(x, y) = 0 \quad \text{Poisson equation}$$

Numerical method:

1. Discretize the derivative:

$$\left(\frac{\partial^2 T}{\partial x^2} \right)_{i,j} \approx \frac{T_{i+1,j} - 2T_{i,j} + T_{i-1,j}}{\Delta x^2}$$

2nd-order, centered scheme

$$\left(\frac{\partial^2 T}{\partial y^2} \right)_{i,j} \approx \frac{T_{i,j+1} - 2T_{i,j} + T_{i,j-1}}{\Delta y^2}$$

$$\Delta x = \Delta y = \Delta$$

With boundary conditions:

$$T(x = 0, y) = L(y), \quad T(x = 1, y) = R(y)$$

$$T(x, y = 0) = D(x), \quad T(x, y = 1) = U(x)$$

2-D (steady) heat equation

Numerical method:

1. Discretize the derivative:

$$\left(\frac{\partial^2 T}{\partial x^2} \right)_{i,j} \approx \frac{T_{i+1,j} - 2T_{i,j} + T_{i-1,j}}{\Delta x^2}$$

2nd-order, centered scheme

$$\left(\frac{\partial^2 T}{\partial y^2} \right)_{i,j} \approx \frac{T_{i,j+1} - 2T_{i,j} + T_{i,j-1}}{\Delta y^2}$$

$$\Delta x = \Delta y = \Delta$$

Final equation:

$$\frac{T_{i+1,j} + T_{i,j+1} - 4T_{i,j} + T_{i-1,j} + T_{i,j-1}}{\Delta^2} = -S_i$$

$$x_i = i * \Delta, \quad i = 0, 1, 2, \dots, N + 1$$

$$y_j = j * \Delta, \quad j = 0, 1, 2, \dots, N + 1$$

**With b.c.'s imposed at $i=0, i=N+1$
and $j=0, j=N+1$**

2-D (steady) heat equation

We now have system of n^2 linear equations, $AT = B$:

$$A = \begin{bmatrix} M & I & 0 & 0 & 0 & \dots & 0 \\ I & M & I & 0 & 0 & \dots & 0 \\ 0 & I & M & I & 0 & \dots & 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots & \vdots & \vdots \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & I & M & I \\ 0 & \dots & 0 & 0 & 0 & I & M \end{bmatrix}$$

- **1D: A was tridiagonal**
- **2D: A is now *block* tridiagonal**
- **M is a $n \times n$ tridiagonal matrix**

$$M = \begin{bmatrix} -4 & 1 & 0 & 0 & 0 & \dots & 0 \\ 1 & -4 & 1 & 0 & 0 & \dots & 0 \\ 0 & 1 & -4 & 1 & 0 & \dots & 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots & \vdots & \vdots \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 1 & -4 & 1 \\ 0 & \dots & 0 & 0 & 0 & 1 & -4 \end{bmatrix}$$

2-D (steady) heat equation

We now have system of n^2 linear equations, $AT = B$:

$$T = \begin{bmatrix} T_{11} \\ T_{21} \\ \vdots \\ T_{n1} \\ T_{12} \\ \vdots \\ T_{n2} \\ \vdots \\ T_{1n} \\ \dots \\ T_{nn} \end{bmatrix} \quad B = \begin{bmatrix} S_{11}\Delta^2 - T_{01} - T_{10} \\ S_{21}\Delta^2 - T_{20} \\ \vdots \\ S_{n1}\Delta^2 - T_{n0} - T_{n+1,1} \\ S_{12}\Delta^2 - T_{02} \\ \vdots \\ S_{n2}\Delta^2 - T_{n+1,2} \\ \vdots \\ S_{n1}\Delta^2 - T_{n+1,1} \\ \vdots \\ S_{nn}\Delta^2 - T_{n+1,n} - T_{n,n+1} \end{bmatrix}$$

Boundary conditions appear in appropriate elements of B

2-D (steady) heat equation

We now have system of n^2 linear equations, $AT = B$

- Direct solution of $n \times n$ matrix: $O(n^3)$ operations (LU decomposition of A + back-substitution with B)
- We have $n^2 \times n^2$ matrix, so on 100×100 grid, matrix is $1e4 \times 1e4$ and contains $1e8$ elements (in double precision, that's 800 mb!)
- So, in 2D (and 3D) direct solution becomes expensive and memory-intensive

2-D (steady) heat equation

We now have system of n^2 linear equations, $AT = B$

- Direct solution of $n \times n$ matrix: $O(n^3)$ operations (LU decomposition of A + back-substitution with B)
- We have $n^2 \times n^2$ matrix, so on 100×100 grid, matrix is $1e4 \times 1e4$ and contains $1e8$ elements (in double precision, that's 800 mb!)
- So, in 2D (and 3D) direct solution becomes expensive and memory-intensive
- But the matrix, A , is *sparse*, so iterative method will only need to store $O(n^2)$ elements (rather than n^4)
- A good iterative method (SOR, conjugate gradient, multigrid) will be faster as well
- Jacobi iteration is *inefficient*, but a good starting point for looking at iterative methods

2-D (steady) heat equation

Jacobi iteration:

$$T_{i,j}^{k+1} = \frac{\Delta^2}{4} S_i + \frac{1}{4} (T_{i+1,j}^k + T_{i,j+1}^k + T_{i-1,j}^k + T_{i,j-1}^k)$$

- **New temperature = Source contribution + average of surrounding temperatures**
- **How do we convert 1D code → 2D?**

2-D (steady) heat equation


Jacobi iteration:

$$T_{i,j}^{k+1} = \frac{\Delta^2}{4} S_i + \frac{1}{4} (T_{i+1,j}^k + T_{i,j+1}^k + T_{i-1,j}^k + T_{i,j-1}^k)$$

- **New temperature = Source contribution + average of surrounding temperatures**
- **How do we convert 1D code → 2D?**
- **Plan:**
 1. **Need 2D variables: x, y, S, T**
 2. **Intialize 2D field and apply boundary conditions**
 3. **During iterations, average in two dimensions rather than one**

2-D (steady) heat equation

Final problem: Use Jacobi iteration to find $T(x,y)$ on unit square with fixed temperature on boundaries and prescribed source, $S(x,y)$



Jacobi iteration in Fortran

1D:

```
do k1=1,kmax
```

```
  Tnew(1:n) = S(1:n)*dx2f + 0.5d0*(T(0:n-1) + T(2:n+1)) !Jacobi
```

```
  deltaT(k1) = maxval(abs(Tnew(1:n)-T(1:n))) !compute relative error
```

```
  T(1:n)=Tnew(1:n)      !update variable
```

```
  if (deltaT(k1)<tol) exit !check convergence criterion
```

```
end do
```


Jacobi iteration in Fortran


2D (see *jacobi2s.f90*):

```
do k1=1, kmax
    Tnew(1:n,1:n) = S(1:n,1:n)*del2f +
0.25*(T(2:n+1,1:n) + T(0:n-1,1:n) + T(1:n,0:n-1) + T(1:n,2:n+1)) !Jacobi
deltaT(k1) = maxval(abs(Tnew(1:n,1:n)-T(1:n,1:n))) !compute relative
error
    T(1:n,1:n)=Tnew(1:n,1:n)      !update variable
    if (deltaT(k1)<tol) exit !check convergence criterion
end do
```

- **Run code with** $U=D=L=0$, $R=1$
- $S = S_0 \sin(\pi x) \sin(\pi y)$

2-D (steady) heat equation


Homogeneous solution ($S=0$):



**The hot wall sets
the temperature
distribution**

2-D (steady) heat equation


Particular solution ($S=10$):



Source in middle
of room shifts
temperature
distribution

2-D (steady) heat equation

Source function:



**Source in middle
of room shifts
temperature
distribution**

2-D (steady) heat equation

Parallelize with OpenMP:

- Essentially the same as in 1D
- Parallel loop iterates across rows of A
- Reduction of ΔT to check for convergence

2-D (steady) heat equation

1D:

```
dmax=0.d0
!$omp parallel do reduction(max:dmax)
do i1=1,n
    Tnew(i1) = S(i1)*dx2f + 0.5d0*(T(i1-1) + T(i1+1))
    dmax = max(dmax,abs(Tnew(i1)-T(i1)))
end do
!$omp end parallel do
deltaT(k1) = dmax
if (deltaT(k1)<tol) exit !check convergence criterion

!$omp parallel do
do i1=1,n
    T(i1) = Tnew(i1)
end do
!$omp end parallel do
```

2-D (steady) heat equation

2D (see *jacobi2s_omp.f90*):

```
dmax = 0.0
!$OMP parallel do reduction(max:dmax)
do j1=1,n
    Tnew(1:n,j1) = S(1:n,j1)*del2f + 0.25*(T(2:n+1,j1) +
    T(0:n-1,j1) + T(1:n,j1-1) + T(1:n,j1+1))!Jacobi iteration
    dmax = max(dmax,maxval(abs(Tnew(1:n,j1)-T(1:n,j1))))
end do
!$OMP end parallel do
deltaT(k1) = dmax
if (deltaT(k1)<tol) exit !check convergence criterion
!$OMP parallel do
do j1=1,n
    T(1:n,j1) = Tnew(1:n,j1)
end do
!$OMP end parallel do
```


2-D (steady) heat equation

- Moving from 1D serial to 2D parallel (with OpenMP) is straightforward
- Much more difficult if solving equations directly
- What about MPI?

2-D (steady) heat equation

What is best domain decomposition?


If we have four processors, can try:



2-D (steady) heat equation

What is best domain decomposition?

If we have four processors, can try:




Then, parallelization is essentially same as differentiation example:

- Loop across rows
- At “boundary” rows, send/recv data needed to compute second derivative
- Reduce $\max(|\Delta T|)$

2-D (steady) heat equation

What is best domain decomposition?

If we have four processors, can try:




- But is the 1D decomposition the best?
- Want to minimize communication
- M “layers”: of a $n \times n$ grid:
 - $(M-1) \cdot n$ boundary points

2-D (steady) heat equation

What is best domain decomposition?

If we have four processors, can also try:




- But is the 1D decomposition the best?
- Want to minimize communication
- M “boxes”: of a $n \times n$ grid:
 - Each interior box has $2n/\sqrt{M}$ boundary points
 - Total: $2n \cdot (\sqrt{M} - 1)$ boundary points

2-D (steady) heat equation

What is best domain decomposition?

If we have four processors, can also try:



- But is the 1D decomposition the best?
- Want to minimize communication
- M “boxes”: of a $n \times n$ grid:
 - Each box has $2n/\sqrt{M}$ boundary points
 - Total: $2n \cdot (\sqrt{M} - 1)$ boundary points
- Boxes: less communication, but more difficult to implement!

2-D (steady) heat equation

- **MPI provides tools for creating and managing complex “topologies”**
- **For example, to create a 4 x 3 “grid” of processes:**

call `MPI_cart_create(MPI_COMM_WORLD, ndims, dims, periods, reorder, new_comm, ierr)`

with: `ndims = 2, dims = (/4,3/), periods = (/false./,true/), reorder = .false.`

- **Here, periods sets periodic boundary conditions along the three columns**

2-D (steady) heat equation

- **MPI provides tools for creating and managing complex “topologies”**
- **For example, to create a 4 x 3 “grid” of processes:**

call `MPI_Cart_create(MPI_COMM_WORLD, ndims, dims, periods, reorder, new_comm, ierr)`

with: `ndims = 2, dims = (/4,3/), periods = (/false./,true/), reorder = .false.`

- **Here, periods sets periodic boundary conditions along the three columns**
- **Other useful commands with the new communicator, new_comm:**
 - **MPI_Cart_coords: given process id, provides (i,j) coordinate**
 - **MPI_Cart_rank: given (i,j), provides id (0, 1, 2, ..., numprocs)**
- **Most useful: MPI_Cart_shift: provides id of neighboring processes in horizontal or vertical direction**
 - **Use to set up send/recv sequences needed for exchanging boundary data.**

2-D (steady) heat equation

- **Most useful:** `MPI_Cart_shift`: provides id of neighboring processes in horizontal or vertical direction
 - Use to set up send/rcv sequences needed for exchanging boundary data.
- How to decide on process grid dimensions?
- `MPI_Dims_create`:
Given number number of processes and dimensions,
outputs process grid
Dimensions (4,3 in picture →)
- e.g. 400 x 300 grid points:
4 x 3 process grid with
100 x 100 points on each grid

