

## Introduction to High Performance Scientific Computing

Autumn, 2016

Lecture 17

Imperial College  
London

Prasun Ray  
5 December 2016

---

---

---

---

---

---

---

## Project Part 1

- Random walks in 1-D
- Consider equations of the form:  

$$X(t+dt) = X(t) + F[X(t), dt]$$
- Here,  $X$  and  $F$ , are both random variables
- Simplest example:  $F = \pm dx$  based on flip of a coin  $\rightarrow$  random walk

Compute several *realizations*,  $X_1(t), X_2(t), \dots, X_M(t)$  and then compute statistics over these  $M$  realizations

These statistics (the mean and stdev) will generally depend on time

Random walk on a network follows same idea

Imperial College  
London

---

---

---

---

---

---


---

## Send/Recv and domain decomposition

A parallel computation computes a potential field,  $f(x, y, z, t)$  on four processors.

$P_0, P_1, P_2, P_3$  solve for  $f$  in separate subdomains

(a)



- How would you compute the gradient?

$$\nabla f = \left( \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z} \right)$$

- No problems with  $x$  and  $z$  directions
- But what about  $y$ ?

Imperial College  
London

---

---

---

---

---


---

---

### Last time: Send/Recv

A parallel computation computes a potential field,  $f(x,y,z,t)$  on four processors.

P0, P1, P2, P3 solve for  $f$  in separate subdomains



Imperial College  
London

- Send and Recv *must* be paired:
    - If yellow sends to blue, blue must recv from yellow
  - Yellow and blue should not send to each other simultaneously
  - Instead:
    - Yellow → Blue
    - Blue → Green
    - Green → Red
    - Red → Yellow (if periodic)
- and then the reverse

---

---

---

---

---

---

---

---

### Parallel differentiation example

```
!Send data at top boundary up to next processor
!i.e. send f(nlocal+1) to myid+1 and store it there as f(1)
!data from myid=numprocs-1 is sent to myid=0
!
if (myid<numprocs-1) then
  receiver = myid+1
else
  receiver = 0
end if

if (myid>0) then
  sender = myid-1
else
  sender = numprocs-1
end if

call MPI_SEND(f(nlocal+1),1,MPI_DOUBLE_PRECISION,receiver,0,
              MPI_COMM_WORLD,ierr)
call MPI_RECV(f(1),1,MPI_DOUBLE_PRECISION,sender,MPI_ANY_TAG,
              MPI_COMM_WORLD,status,ierr)
```

Imperial College  
London

---

---

---

---

---

---

---

---

### Today

Solving the (steady) 2D heat equation

Imperial College  
London

---

---

---

---

---

---

---

---

## Heat equation

Task: Compute temperature distribution in a room

Governing equation: Heat equation (diffusion equation):

$$\frac{\partial T}{\partial t} = \alpha \nabla^2 T + S(\mathbf{x}, t)$$

$$T(\mathbf{x}, t = 0) = f(\mathbf{x}) \quad \text{Initial condition}$$

Here,  $S$  is a *heat source*. Boundary conditions should also be specified as appropriate.

Problem: given the source, initial condition, and boundary conditions, solve for the temperature distribution,  $T(\mathbf{x}, t)$

Imperial College  
London

## 1-D (steady) heat equation

First consider steady problem, e.g.,  $S = S(x)$ ,  $a$  and  $b$  are constants:

$$\frac{\partial^2 T}{\partial x^2} + S(x) = 0 \quad \text{Poisson equation}$$

Numerical method:

1. Discretize the derivative:

$$\frac{\partial^2 T}{\partial x^2} \approx \frac{T_{i+1} - 2T_i + T_{i-1}}{\Delta x^2} \quad \text{2nd-order, centered scheme}$$

$$x_i = i * \Delta x, \quad i = 0, 1, 2, \dots, N + 1$$

$$(N + 1) * \Delta x = 1$$

With boundary conditions:  $T_0 = T_a, T_N = T_b$

Imperial College  
London

## Programming example

$$\text{Equation for } T_i: \frac{T_{i+1} - 2T_i + T_{i-1}}{\Delta x^2} = -S_i$$

In matrix form:  $AT = b$

$$A = \begin{bmatrix} -2 & 1 & 0 & 0 & 0 & \dots & 0 \\ 1 & -2 & 1 & 0 & 0 & \dots & 0 \\ 0 & 1 & -2 & 1 & 0 & \dots & 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \vdots \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 1 & -2 & 1 \\ 0 & \dots & 0 & 0 & 0 & 1 & -2 \end{bmatrix}, \quad b = \frac{1}{\Delta x^2} \begin{bmatrix} -\Delta x^2 T_a - S_1 \\ -S_2 \\ \vdots \\ -S_i \\ \vdots \\ -S_{N-1} \\ -\Delta x^2 T_b - S_N \end{bmatrix}$$

- In 1-D, this is just a tridiagonal system of equations
- Easy to solve directly (with, say, *DGTSV*)

Imperial College  
London

### Jacobi iteration

- Basic idea: rewrite  $Ax=b$  as  $A_1x = A_2x + b$
- Choose  $A_1$  so that it is easy to invert, then solve iterative system:
- $A_1x^{k+1} = A_2x^k + b$ 
  - Requires guess,  $x^0$
- **Jacobi iteration:** Choose  $A_1$  to be diagonal matrix (main diagonal of  $A$ ):

$$\frac{T_{i+1}^{k-1} - 2T_i^k + T_{i-1}^{k-1}}{\Delta x^2} = -S_i$$

$$T_i^k = \frac{\Delta x^2}{2} S_i + \frac{1}{2} (T_{i+1}^{k-1} + T_{i-1}^{k-1})$$

Main algorithm, easy to code!

Imperial College  
London

### Jacobi iteration in Fortran

- One Fortran trick: set variables to be dimension(0:N+1)
- $x(0)=0$ ,  $x(N+1)=1$ ,  $T(0)=a$ ,  $T(N+1)=b$
- Then, easy to compute  $T_1$  using:

$$T_i^k = \frac{\Delta x^2}{2} S_i + \frac{1}{2} (T_{i+1}^{k-1} + T_{i-1}^{k-1})$$

Core part of code (see *jacobi1s.f90*):

```
do k1=1, kmax
  Tnew(1:n) = S(1:n)*dx2f + 0.5d0*(T(0:n-1) + T(2:n+1)) !Jacobi
  deltaT(k1) = maxval(abs(Tnew(1:n)-T(1:n))) !compute relative error
  T(1:n)=Tnew(1:n) !update variable
  if (deltaT(k1)<tol) exit !check convergence criterion
end do
```

Imperial College  
London

### Parallel Jacobi

Let's now parallelize the solver with OpenMP

- Look for loops that can be parallelized
- Look for vectorized operations that can be converted to loops that can be parallelized

Parallel:

```
dmax=0.d0
!$omp parallel do reduction(max:dmax)
do i1=1,n
  Tnew(i1) = S(i1)*dx2f + 0.5d0*(T(i1-1) + T(i1+1))
  dmax = max(dmax,abs(Tnew(i1)-T(i1)))
end do
!$omp end parallel do
deltaT(k1) = dmax
```

Imperial College  
London

### 2-D (steady) heat equation

$$\frac{\partial^2 T}{\partial x^2} + \frac{\partial^2 T}{\partial y^2} + S(x, y) = 0 \quad \text{Poisson equation}$$

Imperial College  
London

---

---

---

---

---

---

---

### 2-D (steady) heat equation

$$\frac{\partial^2 T}{\partial x^2} + \frac{\partial^2 T}{\partial y^2} + S(x, y) = 0 \quad \text{Poisson equation}$$

Numerical method:

1. Discretize the derivative:

$$\left( \frac{\partial^2 T}{\partial x^2} \right)_{i,j} \approx \frac{T_{i+1,j} - 2T_{i,j} + T_{i-1,j}}{\Delta x^2} \quad \text{2nd-order, centered scheme}$$

$$\left( \frac{\partial^2 T}{\partial y^2} \right)_{i,j} \approx \frac{T_{i,j+1} - 2T_{i,j} + T_{i,j-1}}{\Delta y^2}$$

$$\Delta x = \Delta y = \Delta$$

Imperial College  
London

---

---

---

---

---

---

---

### 2-D (steady) heat equation

$$\frac{\partial^2 T}{\partial x^2} + \frac{\partial^2 T}{\partial y^2} + S(x, y) = 0 \quad \text{Poisson equation}$$

Numerical method:

1. Discretize the derivative:

$$\left( \frac{\partial^2 T}{\partial x^2} \right)_{i,j} \approx \frac{T_{i+1,j} - 2T_{i,j} + T_{i-1,j}}{\Delta x^2} \quad \text{2nd-order, centered scheme}$$

$$\left( \frac{\partial^2 T}{\partial y^2} \right)_{i,j} \approx \frac{T_{i,j+1} - 2T_{i,j} + T_{i,j-1}}{\Delta y^2}$$

$$\Delta x = \Delta y = \Delta$$

With boundary conditions:  $T(x=0, y) = L(y)$ ,  $T(x=1, y) = R(y)$   
 $T(x, y=0) = D(x)$ ,  $T(x, y=1) = U(x)$

Imperial College  
London

---

---

---

---

---

---

---

## 2-D (steady) heat equation

Numerical method:

1. Discretize the derivative:

$$\left(\frac{\partial^2 T}{\partial x^2}\right)_{i,j} \approx \frac{T_{i+1,j} - 2T_{i,j} + T_{i-1,j}}{\Delta x^2} \quad \text{2nd-order, centered scheme}$$

$$\left(\frac{\partial^2 T}{\partial y^2}\right)_{i,j} \approx \frac{T_{i,j+1} - 2T_{i,j} + T_{i,j-1}}{\Delta y^2}$$

$$\Delta x = \Delta y = \Delta$$

Final equation:

$$\frac{T_{i+1,j} + T_{i,j+1} - 4T_{i,j} + T_{i-1,j} + T_{i,j-1}}{\Delta^2} = -S_i$$

$$x_i = i * \Delta, \quad i = 0, 1, 2, \dots, N+1$$

$$y_j = j * \Delta, \quad j = 0, 1, 2, \dots, N+1$$

With b.c.'s imposed at  $i=0, i=N+1$   
and  $j=0, j=N+1$

Imperial College  
London

## 2-D (steady) heat equation

We now have system of  $n^2$  linear equations,  $AT = B$ :

$$A = \begin{bmatrix} M & I & 0 & 0 & 0 & \dots & 0 \\ I & M & I & 0 & 0 & \dots & 0 \\ 0 & I & M & I & 0 & \dots & 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & I & M & I \\ 0 & \dots & 0 & 0 & 0 & I & M \end{bmatrix}$$

- 1D: A was tridiagonal
- 2D: A is now *block* tridiagonal
- M is a  $n \times n$  tridiagonal matrix

$$M = \begin{bmatrix} -4 & 1 & 0 & 0 & 0 & \dots & 0 \\ 1 & -4 & 1 & 0 & 0 & \dots & 0 \\ 0 & 1 & -4 & 1 & 0 & \dots & 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 1 & -4 & 1 \\ 0 & \dots & 0 & 0 & 0 & 1 & -4 \end{bmatrix}$$

Imperial College  
London

## 2-D (steady) heat equation

We now have system of  $n^2$  linear equations,  $AT = B$ :

$$T = \begin{bmatrix} T_{11} \\ T_{21} \\ \vdots \\ T_{n1} \\ T_{12} \\ \vdots \\ T_{n2} \\ \vdots \\ T_{1n} \\ \vdots \\ T_{nn} \end{bmatrix} \quad B = \begin{bmatrix} S_{11}\Delta^2 - T_{01} - T_{10} \\ S_{21}\Delta^2 - T_{20} \\ \vdots \\ S_{n1}\Delta^2 - T_{n0} - T_{n+1,1} \\ S_{12}\Delta^2 - T_{02} \\ \vdots \\ S_{n2}\Delta^2 - T_{n+1,2} \\ \vdots \\ S_{n1}\Delta^2 - T_{n+1,1} \\ \vdots \\ S_{nn}\Delta^2 - T_{n+1,n} - T_{n,n+1} \end{bmatrix}$$

Boundary conditions appear in appropriate elements of B

Imperial College  
London

## 2-D (steady) heat equation

We now have system of  $n^2$  linear equations,  $AT = B$

- Direct solution of  $n \times n$  matrix:  $O(n^3)$  operations (LU decomposition of  $A$  + back-substitution with  $B$ )
- We have  $n^2 \times n^2$  matrix, so on  $100 \times 100$  grid, matrix is  $1e4 \times 1e4$  and contains  $1e8$  elements (in double precision, that's 800 mb!)
- So, in 2D (and 3D) direct solution becomes expensive and memory-intensive

Imperial College  
London

---

---

---

---

---

---

---

---

## 2-D (steady) heat equation

We now have system of  $n^2$  linear equations,  $AT = B$

- Direct solution of  $n \times n$  matrix:  $O(n^3)$  operations (LU decomposition of  $A$  + back-substitution with  $B$ )
- We have  $n^2 \times n^2$  matrix, so on  $100 \times 100$  grid, matrix is  $1e4 \times 1e4$  and contains  $1e8$  elements (in double precision, that's 800 mb!)
- So, in 2D (and 3D) direct solution becomes expensive and memory-intensive
- But the matrix,  $A$ , is *sparse*, so iterative method will only need to store  $O(n^2)$  elements (rather than  $n^4$ )
- A good iterative method (SOR, conjugate gradient, multigrid) will be faster as well
- Jacobi iteration is *inefficient*, but a good starting point for looking at iterative methods

Imperial College  
London

---

---

---

---

---

---

---

---

## 2-D (steady) heat equation

Jacobi iteration:

$$T_{i,j}^{k+1} = \frac{\Delta^2}{4} S_i + \frac{1}{4} (T_{i+1,j}^k + T_{i,j+1}^k + T_{i-1,j}^k + T_{i,j-1}^k)$$

- New temperature = Source contribution + average of surrounding temperatures
- How do we convert 1D code  $\rightarrow$  2D?

Imperial College  
London

---

---

---

---

---

---

---

---

## 2-D (steady) heat equation

Jacobi iteration:

$$T_{i,j}^{k+1} = \frac{\Delta^2}{4} S_i + \frac{1}{4} (T_{i+1,j}^k + T_{i,j+1}^k + T_{i-1,j}^k + T_{i,j-1}^k)$$

- New temperature = Source contribution + average of surrounding temperatures
- How do we convert 1D code → 2D?
- Plan:
  1. Need 2D variables: x, y, S, T
  2. Initialize 2D field and apply boundary conditions
  3. During iterations, average in two dimensions rather than one

Imperial College  
London

---

---

---

---

---


---

---

---

## 2-D (steady) heat equation

Final problem: Use Jacobi iteration to find  $T(x,y)$  on unit square with fixed temperature on boundaries and prescribed source,  $S(x,y)$



Imperial College  
London

---

---

---

---

---

---

---

---

## Jacobi iteration in Fortran

```
1D:
do k=1, kmax
  Tnew(1:n) = S(1:n)*dx2f + 0.5d0*(T(0:n-1) + T(2:n+1)) !Jacobi
  deltaT(k1) = maxval(abs(Tnew(1:n)-T(1:n))) !compute relative error
  T(1:n)=Tnew(1:n) !update variable
  if (deltaT(k1)<tol) exit !check convergence criterion
end do
```

Imperial College  
London

---

---

---

---

---

---

---

---



### Jacobi iteration in Fortran

2D (see *jacobi2s.f90*):

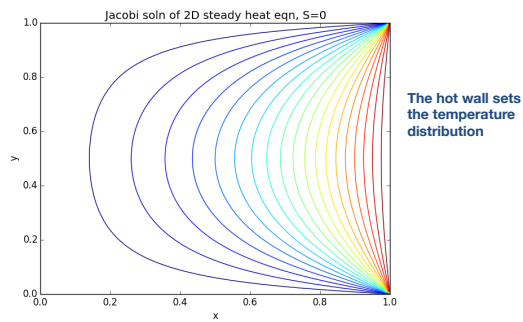
```
do k1=1,kmax
  Tnew(1:n,1:n) = S(1:n,1:n)*del2f +
0.25*(T(2:n+1,1:n) + T(0:n-1,1:n) + T(1:n,0:n-1) + T(1:n,2:n+1)) !Jacobi
  deltaT(k1) = maxval(abs(Tnew(1:n,1:n)-T(1:n,1:n))) !compute relative
error
  T(1:n,1:n)=Tnew(1:n,1:n) !update variable
  if (deltaT(k1)<tol) exit !check convergence criterion
end do
```

- Run code with  $U=D=L=0$ ,  $R=1$
- $S = S_0 \sin(\pi x) \sin(\pi y)$

Imperial College  
London

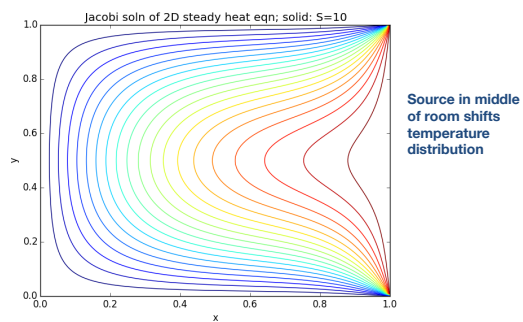
### 2-D (steady) heat equation

Homogeneous solution ( $S=0$ ):




### 2-D (steady) heat equation

Particular solution ( $S=10$ ):



## 2-D (steady) heat equation

Source function:




---

---

---

---

---

---

---

---

## 2-D (steady) heat equation

Parallelize with OpenMP:

- Essentially the same as in 1D
- Parallel loop iterates across rows of A
- Reduction of deltaT to check for convergence

Imperial College  
London

---

---

---

---

---

---

---

---

## 2-D (steady) heat equation

1D:

```
dmax=0.d0
!$omp parallel do reduction(max:dmax)
do i1=1,n
  Tnew(i1) = S(i1)*dx2f + 0.5d0*(T(i1-1) + T(i1+1))
  dmax = max(dmax,abs(Tnew(i1)-T(i1)))
end do
!$omp end parallel do
deltaT(k1) = dmax
if (deltaT(k1)<tol) exit !check convergence criterion
!$omp parallel do
do i1=1,n
  T(i1) = Tnew(i1)
end do
!$omp end parallel do
```

Imperial College  
London

---

---

---

---

---

---

---

---

## 2-D (steady) heat equation

2D (see `jacobi2s_omp.f90`):

```
dmax = 0.0
!$OMP parallel do reduction(max:dmax)
do j1=1,n
  Tnew(1:n,j1) = S(1:n,j1)*del2f + 0.25*(T(2:n+1,j1) +
    T(0:n-1,j1) + T(1:n,j1-1) + T(1:n,j1+1))!Jacobi iteration
  dmax = max(dmax,maxval(abs(Tnew(1:n,j1)-T(1:n,j1))))
end do
!$OMP end parallel do
deltaT(k1) = dmax
if (deltaT(k1)<tol) exit !check convergence criterion
!$OMP parallel do
do j1=1,n
  T(1:n,j1) = Tnew(1:n,j1)
end do
!$OMP end parallel do
```

Imperial College  
London

---

---

---

---

---

---

---

---

## 2-D (steady) heat equation

- Moving from 1D serial to 2D parallel (with OpenMP) is straightforward
- Much more difficult if solving equations directly
- What about MPI?

Imperial College  
London

---

---

---

---

---

---


---

---

## 2-D (steady) heat equation

What is best domain decomposition?

If we have four processors, can try:



Imperial College  
London

---

---

---

---

---

---


---

---

## 2-D (steady) heat equation

What is best domain decomposition?

If we have four processors, can try:



Then, parallelization is essentially same as differentiation example:

- Loop across rows
- At "boundary" rows, send/recv data needed to compute second derivative
- Reduce  $\max(|\Delta T|)$

Imperial College  
London

---

---

---

---

---

---


---

---

## 2-D (steady) heat equation

What is best domain decomposition?

If we have four processors, can try:



- But is the 1D decomposition the best?
- Want to minimize communication
- M "layers": of a  $n \times n$  grid:
  - $(M-1)n$  boundary points

Imperial College  
London

---

---

---

---

---

---


---

---

## 2-D (steady) heat equation

What is best domain decomposition?

If we have four processors, can also try:



- But is the 1D decomposition the best?
- Want to minimize communication
- M "boxes": of a  $n \times n$  grid:
  - Each interior box has  $2n/\sqrt{M}$  boundary points
  - Total:  $2n(\sqrt{M}-1)$  boundary points

Imperial College  
London

---

---

---

---

---

---


---

---

## 2-D (steady) heat equation

What is best domain decomposition?

If we have four processors, can also try:



- But is the 1D decomposition the best?
- Want to minimize communication
- M "boxes": of a  $n \times n$  grid:
  - Each box has  $2n/\sqrt{M}$  boundary points
  - Total:  $2n(\sqrt{M} - 1)$  boundary points
- Boxes: less communication, but more difficult to implement!

Imperial College  
London

---

---

---

---

---

---

---

---

## 2-D (steady) heat equation

- MPI provides tools for creating and managing complex "topologies"
- For example, to create a  $4 \times 3$  "grid" of processes:  
 call `MPI_cart_create(MPI_COMM_WORLD, ndims, dims, periods, reorder, new_comm, ierr)`  
 with: `ndims = 2, dims = (/4,3/), periods = (/false,/.true/), reorder = .false.`
- Here, `periods` sets periodic boundary conditions along the three columns

Imperial College  
London

---

---

---

---

---

---

---

---

## 2-D (steady) heat equation

- MPI provides tools for creating and managing complex "topologies"
- For example, to create a  $4 \times 3$  "grid" of processes:  
 call `MPI_Cart_create(MPI_COMM_WORLD, ndims, dims, periods, reorder, new_comm, ierr)`  
 with: `ndims = 2, dims = (/4,3/), periods = (/false,/.true/), reorder = .false.`
- Here, `periods` sets periodic boundary conditions along the three columns
- Other useful commands with the new communicator, `new_comm`:
  - `MPI_Cart_coords`: given process id, provides (i,j) coordinate
  - `MPI_Cart_rank`: given (i,j), provides id (0, 1, 2, ..., numprocs)
- Most useful: `MPI_Cart_shift`: provides id of neighboring processes in horizontal or vertical direction
  - Use to set up send/rcv sequences needed for exchanging boundary data.

Imperial College  
London

---

---

---

---

---


---

---

---

## 2-D (steady) heat equation

- Most useful: `MPI_Cart_shift`: provides id of neighboring processes in horizontal or vertical direction
  - Use to set up send/rcv sequences needed for exchanging boundary data.
- How to decide on process grid dimensions?
- `MPI_Dims_create`:  
Given number number of processes and dimensions, outputs process grid  
Dimensions (4,3 in picture →)
- e.g. 400 x 300 grid points:  
4 x 3 process grid with  
100 x 100 points on each grid



Imperial College  
London

---

---

---

---

---

---

---

---